




# Proof checking in the Linux kernel

Jacob Herbst (mwr148)  
Institute of Computer Science (DIKU)



# Agenda

- eBPF, Proof Carrying Code and why we need a proof-checker in the kernel
- Logical Framework with Side Conditions (LFSC)
- An LF proof
- Shortcomings of the implementation
- Conclusion



# eBPF and why we want formal verification

- eBPF is a virtual machine that gives sandbox functionality with kernel privileges.
- limbo between usability and safety
- programs are checked by static analysis called the eBPF verifier.
  - 19k lines of code.
  - many bugs, making eBPF security hazard
  - disabled by default in most Linux distributions.




Figure: Loading process of eBPF<sup>1</sup>

<sup>1</sup><https://ebpf.io/what-is-ebpf/>

Jacob Herbst (mwr148) (DIKU) — Proof checking in the Linux kernel



# Proof Carrying Code

- code producer is responsible for ensuring the correctness
- code consumer shall check that proofs have not been tampered and that a certificate is valid.

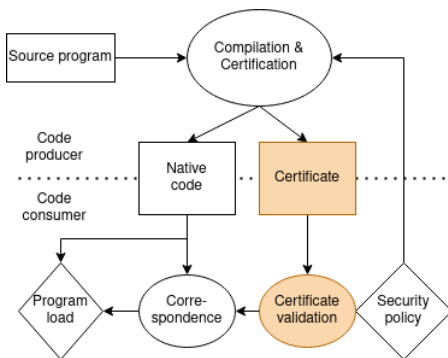


Figure: Structure of PCC

# Logical Framework - with Side Conditions?

- An extension of the simply typed lambda calculus with dependent product type
  - $\prod_{x:A} B(x)$ , with  $A : \mathcal{U}$  and  $B : A \rightarrow \mathcal{U}$
  - if  $x \notin FV(B)$  then  $A \rightarrow B$
- Used for computer-assisted proofs using the Curry-Howard correspondence

Formal Logic	Type Theory
$\top$ (true)	$()$
$\perp$ (false)	void
$\wedge$ (conjunction)	product type
$\vee$ (disjunction)	sum type
$\rightarrow$ (implication)	function type
$\forall$ (universal quantification)	$\Pi$ (dependent type)

- LFI extend LF with implicit arguments
- LFSC extend LFI with Side Conditions



# LFSC syntax

$$\begin{aligned}
 K &::= \text{type} \mid \text{type}^c \mid \text{kind} \mid \Pi x : A. K \mid \text{int} \mid \text{rational} \\
 A, B &::= a \mid A M \mid \Pi x : \{S M\}. A \mid \Pi x : A. B \\
 M, N, O &::= x \mid c \mid z \mid q \mid * \mid M : A \mid \text{let } x M N \mid \lambda x. M \mid \lambda x : A. M \mid M N \\
 P &::= c \mid c x_1 \dots x_n \\
 S, T, U &::= x \mid c \mid -S \mid S \oplus S \mid c S_1 \dots S_n \mid \text{let } x S T \mid \text{markvar } S \mid \\
 &\quad \text{ifequal } S_1 S_2 T U \mid \text{match } S (P_1 T_1) \dots (P_n T_n) \mid \text{fail } S \mid \\
 &\quad \text{ifneg } S T U \mid \text{ifzero } S T U \mid \text{ifmarked } S T U \mid \text{ztoq } S \\
 \oplus &\in \{+, /, *\}
 \end{aligned}$$

Figure: Syntactical categories of LFSC



# LFSC typing

$$\begin{array}{c}
 \frac{\vdash \Gamma}{\Gamma \vdash \mathbf{type} \Rightarrow \mathbf{kind}} \quad
 \frac{\vdash \Gamma}{\Gamma \vdash \mathbf{int} \Rightarrow \mathbf{type}} \quad
 \frac{\vdash \Gamma}{\Gamma \vdash \mathbf{rational} \Rightarrow \mathbf{type}} \\
 \frac{\vdash \Gamma \quad x : A \in \Gamma}{\Gamma \vdash x \Rightarrow A} \quad
 \frac{\vdash \Gamma \quad a : K \in \Sigma}{\Gamma \vdash a \Rightarrow K} \quad
 \frac{\vdash \Gamma \quad c : A \in \Sigma}{\Gamma \vdash c \Rightarrow A} \\
 \frac{\Gamma \vdash_{\Sigma} A \Leftarrow \mathbf{type} \quad \Gamma, x : A \vdash_{\Sigma} C \Rightarrow \alpha \quad \alpha \in \{\mathbf{type}, \mathbf{type}^c, \mathbf{kind}\}}{\Gamma \vdash_{\Sigma} \Pi x : A. C \Rightarrow \alpha} \\
 \\
 \frac{\Gamma \vdash_{\Sigma} A \Rightarrow \Pi x : B. K \quad \Gamma \vdash_{\Sigma} M \Leftarrow B}{\Gamma \vdash_{\Sigma} A M \Rightarrow [M/x]K} \quad
 \frac{\Gamma \vdash_{\Sigma} M \Rightarrow \Pi x : A. B}{\Gamma \vdash_{\Sigma} M * \Rightarrow [* / x]B} \\
 \\
 \frac{\Gamma \vdash_{\Sigma} M \Rightarrow \Pi x : A. B \quad \Gamma \vdash_{\Sigma} N \Leftarrow A}{\Gamma \vdash_{\Sigma} M N \Rightarrow [N/x]B} \quad
 \frac{\Gamma \vdash_{\Sigma} M \Leftarrow A}{\Gamma \vdash_{\Sigma} M : A \Rightarrow A} \\
 \\
 \frac{\Gamma \vdash_{\Sigma} A \Rightarrow \mathbf{type} \quad \Gamma, x : A \vdash_{\Sigma} M \Rightarrow B}{\Gamma \vdash_{\Sigma} \lambda x : A. M \Rightarrow \Pi x : A. B} \quad
 \frac{\vdash \Gamma}{\Gamma \vdash_{\Sigma} q \Rightarrow \mathbf{rational}} \\
 \\
 \frac{\Gamma, x : A \vdash_{\Sigma} M \Rightarrow B}{\Gamma \vdash_{\Sigma} \lambda x. M \Leftarrow \Pi x : A. B} \quad
 \frac{\vdash \Gamma}{\Gamma \vdash z \Rightarrow \mathbf{integer}}
 \end{array}$$

Figure: Bidirectional typing rules for LFSC



# 2 + 2 = 4 using cvc5 and LFSC

```

1 (check
2 (@ t1 (int 4)
3 (@ t2 (int 2)
4 (@ t3 (= (a.+ t2 (a.+ t2 (int 0))) t1)
5 (# a0 (holds (not t3))
6 (: (holds false)
7 (eq_resolve _ _ a0
8 (trans _ _ _
9 (cong _ _ _ _
10 (refl f_not)
11 (trans _ _ _
12 (cong _ _ _ _
13 (cong _ _ _ _
14 (refl f_=)
15 (trust t3))
16 (refl t1))
17 (trust (= (= t1 t1) true))))
18 (trust (= (not true) false))))))

```

$$\text{eq\_resolve} : \prod_{f:T} \prod_{g:T} \prod_{p_1:[f]} \prod_{p_2:[f=g]} [g]$$

$$\text{trans} : \prod_{t_1:T} \prod_{t_2:T} \prod_{t_3:T} \prod_{u_1:[t_1=t_2]} \prod_{u_2:[t_1=t_3]} [t_1=t_3]$$

$$\text{cong} : \prod_{a_1:T} \prod_{b_1:T} \prod_{a_2:T} \prod_{b_2:T} \prod_{u_1:[a_1=b_1]} \prod_{u_2:[a_2=b_2]} [a_1 a_2 = b_1 b_2]$$

$$\text{refl} : \prod_{t:T} [t=t]$$

$$\text{holds} : \prod_{t:T} T$$

$$= \stackrel{\text{def}}{=} \lambda t_1 : T. \lambda t_2 : T. f_{=t_1} t_1$$

$$\begin{array}{c}
\frac{f_=}{(=) = (=)} \quad \text{refl (14)} \quad \frac{t_3}{(2 + 2 = 4)} \quad \text{trust (15)} \\
\hline
(2 + 2 =) = (4 =) \quad \text{cong (13)} \quad \frac{t_1}{4 = 4} \quad \text{refl (16)} \quad (4 = 4) = \top \\
\hline
(2 + 2 = 4) = (4 = 4) \quad \text{cong (12)} \quad (4 = 4) = \top \quad \text{trust (17)} \\
\hline
(2 + 2 = 4) = \top \quad \text{trans (11)} \\
\hline
(2 + 2 = 4) = \top
\end{array}$$

$$\begin{array}{c}
\frac{\neg}{\neg \neg} \quad \text{refl (10)} \quad \frac{\neg(2 + 2 = 4)}{\neg(2 + 2 = 4)} \quad a_0(5) \\
\hline
\neg(2 + 2 = 4) \\
\hline
\frac{\neg(2 + 2 = 4) = \neg \top}{\neg(2 + 2 = 4) = \perp} \quad \text{cong (9)} \quad \frac{\neg \top = \perp}{\neg \top = \perp} \quad \text{trust (18)} \\
\hline
\neg(2 + 2 = 4) = \perp \quad \text{eq\_resolve (7)} \quad \text{trans (8)}
\end{array}$$





# Demo



# Demo

Why can we also prove  $2 + 3 = 4$ ?



# Demo

Why can we also prove  $2 + 3 = 4$ ?

```
1 (declare apply (! t1 term (! t2 term term)))  
2 (declare f_a.+ term)  
3 (define a.+ (# x term (# y term (apply (apply f_a.+ x) y))))
```

$$a.+ : (\lambda x : term. \lambda y : term. term)$$


# Demo

Why can we also prove  $2 + 3 = 4$ ?

```
1 (declare apply (! t1 term (! t2 term term)))
2
3 (declare f_a.+ term)
4
5 (define a.+ (# x term (# y term (apply (apply f_a.+ x) y))))
```

$$a.+ : (\lambda x : term. \lambda y : term. term)$$

We could just as well use  $-$ , as the definition is as follows:


```
1 (declare f_a.- term)
2 (define a.- (# x term (# y term (apply (apply f_a.- x) y))))
```

This is a problem if we want to use it to make proofs by hand, but in a PCC context, it will suffice, as long as we are sure the in the kernel VC generator is sound.



# Shortcomings

- Implementation in this thesis is not fast.
- Approach follows a typically staged compilation. Parser → Transformation → Typechecking.
- Overhead in both memory and runtime.
- *Ifscr* takes an online approach, where parsing, and type checking happens all at once (no intermediate stages)



# Conclusion

- Side Conditions might be unnecessary
- Formal verification of eBPF programs will take longer than the eBPF verifier but can be done fairly efficiently, but not with the approach tried in my work.
- Proposal is to take a similar approach to *lfsc* and implement it in Rust, as it seems like a doable solution.

