



II

Msc. Thesis

Jacob Herbst (mwr148)

Verifying eBPF programs in the Linux Kernel

Investigation in feasibility of checking Verification conditions for as part of the eBPF syscall

Repository: <https://github.com/Spatenheinz/LeSpeciale>

Advisor: Ken Friis Larsen

2023-05-31

Abstract

eBPF is a Linux kernel sub-system that allows users to run code of a limited assembly-like language inside the Linux kernel. This can potentially be dangerous in the hands of malicious or even uncaring users. To combat this, the Linux kernel provides static analysis of eBPF programs to reject harmful programs. This static analysis has in the past shown to be unsound, meaning harmful programs have been accepted. In this project, we investigate the feasibility of making a logical proof checker that runs inside the kernel, implemented in Rust. We do so by considering the Logical Framework with Side Conditions (LFSC) language. The motivation behind such a proof checker is to be able to formally prove the correctness of eBPF programs. We consider the feasibility not only as a standalone tool but as a hypothetical part of a larger Proof Carrying Code (PCC) architecture. By this, we provide an analysis of how eBPF programs are loaded and what the static analysis consists of and hereby describe a vision for a PCC architecture. The main focus of the report is to describe and evaluate the implementation of LFSC with respect to eBPF, PCC, and Rust in the Linux kernel. Although the implementation is not as fast as other LFSC proof checkers it does show decent promise for a part of a PCC system for eBPF.

Contents

1	Introduction	4
2	Background	5
2.1	Linux and the eBPF subsystem	5
2.2	Proof Carrying Code	7
2.3	Rust	7
2.4	Curry Howard Isomorphism	11
3	Proof Carrying Code in the Kernel	12
3.1	Spelunking the eBPF verifier	12
3.2	eBPF and PCC	15
4	Logical Framework with Side Conditions	17
4.1	Abstract Syntax	18
4.2	Signatures and Contexts	18
4.3	Typing	19
4.4	Operational Semantics of Side conditions	22
4.5	Concrete Syntax	24
4.6	Using LFSC to show $P \wedge \neg P$ is unsatisfiable	24
5	The in-kernel proof checker - LFSCR	27
5.1	General implementation choices	27
5.2	Reading and formatting proofs	29
5.3	Typechecking LFSC	33
6	Experiments	41
7	Evaluation	43
7.1	Speed	43
7.2	Memory	46

7.3 LFSCR - strong suits and weaknesses	47
8 Is PCC a good idea?	48
9 Future work	49
10 Conclusion	49

1 Introduction

eBPF is a sub-system in the Linux kernel, which allows users of the system to dynamically load eBPF bytecode into the kernel. The program can then be executed when certain events happen. This technology enables many interesting features, such as high-speed package filtering (which was the initial intent of the system) and Express Data Path (XDP) by circumventing the network stack of the operating system. Furthermore, it can be used for system monitoring by access to kernel probes, etc.

eBPF allows untrusted users to run arbitrary code in the kernel, which in itself is no problem if the program is non-malicious, but can be detrimental if not. Programs that run in the Linux kernel must therefore be both safe and secure. An unsafe program in the Linux kernel might break the system altogether, whilst a malicious user could get access to confidential information. To prevent this the eBPF sub-system will perform an abstract interpretation of a program and then reject unsafe programs, before they are loaded into the kernel. This process is called the verifier. The verifier has been subject to multiple bugs in the past, which can lead to privilege escalation[15][16]. On the other hand, the verifier also limits the capabilities of eBPF programs, by outright rejecting programs with loops. This is a design choice since non-halting programs should not live in the kernel, and there is no way to tell if a program will halt using static analysis, as per the halting problem. The result is a security mechanism that is both overly conservative in some cases and unsound.

An alternative to static analysis is to verify the safety of programs using formal logic. This can be done by generating a logical formula that describes the properties of a program, a so-called verification condition, and then checking the unsatisfiability of the negated formula, with an SMT solver, which makes the formula valid. The process of generating a verification condition is cheap, however, checking the validity of the formula can be a computationally heavy task. Proof Carrying Code is an architecture and security mechanism first introduced by Necula in 1998[14]. Here the user is responsible for proving that a certain program follows a set of security parameters and then the kernel only has to check this proof. The general concept is described in detail in Section 2.2. All in all these tasks are far cheaper than producing a proof.

In this project, we investigate the feasibility of a proof checker that can run inside the Linux kernel. The implementation not only tests this subcomponent of proof-carrying code but also tests if the Rust programming language, introduced in the Linux kernel 6.1 and forward, is a good tool for the job. The proof checker leverages the Curry-Howard correspondence and thus checking a proof amounts to type-checking. We specifically consider the dependently typed meta-framework, Logical Framework with Side Conditions (LFSC)[17]. We describe the semantics of the type checking and present an implementation.

The report is structured as follows. In Section 2 we give a brief introduction to eBPF as a sub-system in Linux, Proof Carrying Code, Rust, and its limitations. Then in Section 3 we go into detail about how the eBPF verifier works and discuss Proof Carrying Code in the context of eBPF. We then present the syntax and semantics of the LFSC in Section 4, to serve as a basis for the implementation of the proof-checker. We do not prove any properties of the type system. In Section 5 we present the design choices and implementation details of the type-checker called LFSCR. We then evaluate the performance and general usability of LFSCR in a PCC system Section 7, by considering proofs generated from a simple verification condition generator, presented in Section 6. Lastly, we conclude the project and present some ideas for future work in 8 and 10.

2 Background

2.1 Linux and the eBPF subsystem

The Linux kernel is a monolithic kernel, meaning that it provides not only the fundamental services necessary for a fully functioning operating system, but also a virtual interface for communication with hardware, including filesystems, network stacks, and device drivers, among others, all in a single executable. This contrast with micro kernels, where only the core functionality of the operating system exists in kernel and all other functionality lives in userspace. In both cases the standard form of communication between processes in userspace and the kernel is through system calls. This for instance happens when accessing hardware, requesting memory etc.

One major advantage of monolithic kernels is the minimal overhead of system level tasks as there is no need for interprocess communication from the different parts of the system. In a micro kernels parts of the functionality provided by a monolithic kernel live in userspace, such as filesystems. Since processes does not have direct “knowledge” of other processes, there will be an overhead when doing some system level tasks. On the other hand monolithic kernels in general provides little to no extensionality, where microkernels on the other hand is design for exactly this.

Albeit being a monolithic architecture the Linux kernel is characterized by its modular nature. The modularity is achieved either with dynamically loadable kernel modules (LKMs) or through the eBPF subsystem. LKMs are more general and can extend functionality and support a diverse range of services, such as device drivers, virtual filesystems etc. They function similarly to traditional filesystems in their execution and can be loaded and unloaded as necessary. They provide great power but at the same time no safety guarantees. Thus LKMs have inherent security risks and require root privileges, limiting their use to trusted users. A malicious LKM can destroy a kernel completely, especially considering that kernel modules may be proprietary. This creates a tradeoff between extensionality and trust.

eBPF on the other hand is more limited in capabilities but provides a platform for doing smaller tasks which still leverages the power of the kernel without the same level of risk.

2.1.1 eBPF

BPF (Berkeley Package Filter) as it was originally presented is a system for effective filtering of network packages by allowing dynamical loading of package filters from userspace into the network stack in kernel space. eBPF, an extension to BPF, was later introduced in the Linux kernel, offering a different approach to kernel extension. Since then many other eBPF’s have appeared but we only consider the eBPF in the Linux kernel. The eBPF virtual machine leverages the privileges of the kernel to oversee the entire system, enabling more powerful control. The eBPF name furthermore refers to a just-in-time (JIT) compiled reduced instruction set computer (RISC). The ecosystem seeks to ensure security by performing static analysis of the limited language. Whilst the language itself is turing complete, the static analysis explicitly disallow backwards jumps. In eBPF there are 11 registers, r0-r10, where r0 is used for return values of functions and r10 is a readonly framepointer to a 512 byte stack. Instructions can be moves, addition and jumps in a similar manner to other reduced instructions sets.

The eBPF program loading process is presented in Figure 1. First step in the process involves obtaining a program using an abstraction tool such as BCC[3] or libbpf[12] or writing the program by hand in C macros. The `bpf()` syscall is then invoked and the eBPF bytecode is moved into the kernel. The verifier then performs a series of security measures that determines if a program is accepted for loading or rejected. The security measures involves static analysis in the form of abstract interpretation using tristate numbers, cycle detection, division by zero, and more. We describe this in detail in Section 3.1. If a program is determined safe by the verifier, the program is loaded into the kernel, either for interpretation or JIT-compiled depending on the kernel configurations. The program is then attached to a hook.




Figure 1: eBPF loading process [13]

eBPF programs are event-driven, meaning they can be attached to a certain hook, and every time the corresponding event occurs, the program is triggered. For example, a program can be attached to a socket, and every time something is written to this socket, the program is triggered. Although an eBPF program lives in kernel space, it conceptually resides somewhere between user and kernel space. It can interact with both the kernel and user space through a collection of key-value stores called maps, realized as a variety of different data structures such as ring buffers, arrays, etc. These data structures also reside in kernel space and are constructed through the `bpf` syscall, allowing eBPF programs to read and write to a map. Users can also read and write to the maps using the `bpf` syscall, thus serving as a communication layer between eBPF program and user space.

eBPF programs should likewise not be seen as a coherent part of the kernel as it cannot call kernel functions directly. The reason for this is that eBPF is designed to be kernel version agnostic[7], and Linux kernel functions are not stable. In reality it is not agnostic since different versions of the verifier will reject and accept different programs. Instead, the eBPF subsystem provides a stable API of helper functions to provide functionality not immediately accessible in the limited instruction set. Furthermore it is possible to have eBPF programs call other eBPF programs, and even chain them together in tail calls similar to standard function calls, greatly extending the functionality of eBPF programs.

This subsystem hence serves an opportunity for users to extend the kernel with functionality that can run in a sandboxed manner while still leveraging the powers of the kernel. Unfortunately due to bugs in the verifier, most major Linux distributions such as Ubuntu, Fedora, Redhat and many more disallow non-root users to load eBPF programs.

2.2 Proof Carrying Code

We are interested in investigating Proof Carrying Code as an alternative to the eBPF verifier. Proof Carrying Code (PCC) is a mechanism designed to ensure the safe execution of programs from untrusted sources. The PCC process is presented in Figure 2. In this project we specifically consider the marked parts of the process.

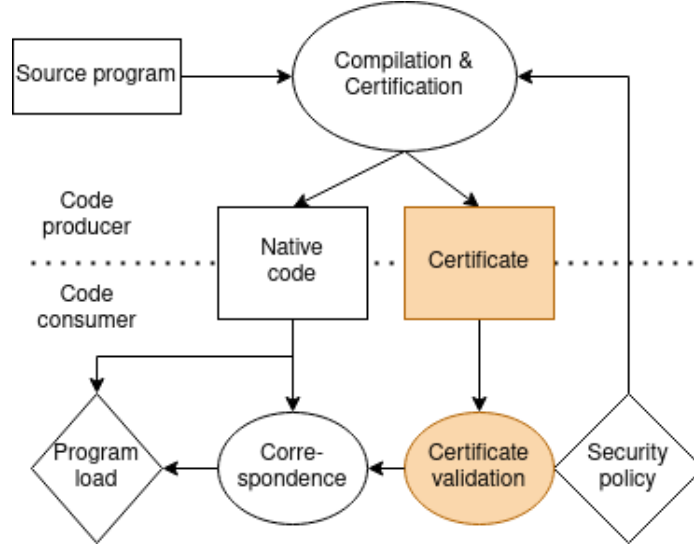


Figure 2: Structure of Proof Carrying Code

The PCC architecture can be divided into two spaces, namely the code producer and the code consumer. The code producer aims to execute some code at the expense of the code consumer. In case there is not perfect trust between code consumer and code producer, the code consumer must protect itself from potentially malicious or unsafe programs. To achieve this, the code consumer issues a collection of safety rules, referred to as the safety policy, which the code producer must adhere to. Depending on the specific domain, the safety policy could include constraints such as loop termination, no out-of-bounds memory operations along with other constraints that ensures the code consumer is not getting corrupted. The code producer then uses the safety policy to certify the program in a compilation stage. The certification process produces both the native code to be executed at the consumer and the certificate for the safety of the native code. This process can potentially be computationally heavy, and it is preferable for the code producer to perform the certification and compilation.

The next stage is the verification phase, where the producer hands over the results of the certification process to the consumer. The consumer then checks the validity of the proof in two ways. First, the safety proof must be valid modulo the safety policies, and second, the safety proof must correspond to the native code. This process should be quick and follow an algorithm trusted by the consumer. If both criteria are met, the consumer can mark the native code as safe and proceed to execute the program, possibly multiple times.

2.3 Rust

In this section, we give a short description of the Rust programming language to motivate why it is used as host language in this project. Rust is a modern systems programming language that was first introduced in 2010. It was designed to address common issues faced by developers in

writing low-level, high-performance code, such as memory safety and thread safety.

One of the key features of Rust is its ownership model, which ensures that memory is managed efficiently and safely. Rust's ownership model is based on the concept of ownership and borrowing, which allows the compiler to track the lifetime of objects and manage memory without the need for a garbage collector. This ensures that common issues such as null pointer dereferences, use-after-free errors, and buffer overflows can occur, while at the same time being comparable to C in performance.

Rust's syntax is similar to that of C and C++, but it also includes modern language features like pattern matching, closures, and iterators. Rust also has a strong focus on performance and optimization, which makes it an ideal language for building high-performance applications and systems.

An intrinsic part of the Rust language is the borrow checker; a form of static analysis which ensures that a program complies with the ownership rules. The rules are 3 fold.

- Each value has an owner.
- There can only be one owner for each value.
- When the owner goes out of scope, its values are freed (dropped in Rust terminology).

If we for instance consider:

```
fn main() {  
    let x : i32 = 42;  
}
```

Then the value 42 has an owner `x`. When the `main()` function ends then the owner `x` goes out of scope and the value is dropped. This specific type of `x` is `i32` and will thus reside on the stack, however, if we needed something that was heap allocated then we could create a value using a `Box`:

```
fn main() {  
    let y : Box<i32> = Box::new(42);  
}
```

Then `y` will be a `Box` type, which is the simplest form of heap allocation. Objects in Rust can implement a trait, essentially an interface, called `Drop`, and define a `drop` function. When `y` goes out of scope a call to `drop` happens and the memory will be freed.

Ownership can be transferred by either `move`, `copy`, or `clone`. For primitive types, those that usually reside on the stack, they can be copied from one variable to another. Heap-allocated values can be either cloned, which essentially works as a `memcpy`, or by moving it. This means that the snippet below will get a compile-time error at line 6 because the variable that owns the box containing 42 is no longer owned by `y` but rather by `a`.

```
fn main() {  
    let y : Box<i32> = Box::new(42);  
    let z = y.clone();  
    assert_eq!(y, z);  
    let a = y;  
    assert_eq!(y, z);  
}
```

Since cloning is a `memcpy` it can be quite inefficient and should most often be avoided. Rust allows pass-by-reference in the form of borrowing. Borrowing describes the action of receiving something with the promise of return. When borrowing a value the memory address of the value is referenced by an `&` and is essentially just a pointer. It is worth noting that a reference differs from pointers in C in that they cannot be `NULL` and thus by mere construction eliminates `NULL` pointer problems. There are two types of references, exclusive and shared references. Exclusive references can mutate the borrowed value, while a shared reference may only read the reference. There are 3 rules that borrowing is subject to:

1. There can exist either a single exclusive reference or multiple shared references at a given time.
2. References must always be valid, which means it is impossible to borrow a value after its owner has gone out of scope.
3. A value cannot be modified whilst referenced.

These rules invariantly ensure that a reference is always as perceived to the borrower. If multiple mutable borrows were allowed at the same time, or even simultaneously with a shared reference, then one of the mutable borrows may destroy the reference for the others. An example where this is extremely obvious is when we consider a reference to a data structure that might need to be reallocated, such as dynamic arrays. In such a situation the address of the old array will no longer be valid. This is ensured to never happen according to rule 3.

Rust also promises zero-cost abstractions and high-level features such as sum types, pattern matching, and traits/interfaces, while still having performance similar to C.

These promises of memory safety and high-level abstractions are what constitute the basics of the Rust programming language. Safe memory management without the overhead of a garbage collector has gained Rust popularity in recent years, especially in the systems programming community, due to its combination of performance, safety, and ease of use. This popularity also includes the Linux Development community, where safety, security, and reliability are mission-critical. As of kernel 6.1 Rust is officially supported in the Linux kernel, albeit fairly limited as described throughout this report.

Hence Rust seems like an appropriate tool for a program that has to run in the kernel, and where robustness is critical, but we must also consider what restrictions the kernel imposes on Rust.

2.3.1 Rust in the kernel?

As of kernel version 6.2.8rc, the Rust kernel development framework not a lot of functionality is exposed. The crates/modules immediately exposed in the kernel are `alloc`, `core`, `kernel`, `compiler_builtins`, and `macros`. The `macros` crate is tiny and exposes the ability to easily describe an LKM's meta-data. The `compiler_builtins` are compiler built-in functionality that usually resides in the standard library `std`. The builtins supported in the kernel at the moment are nothing more than panics (exceptions). The `kernel` crate exposes the kernel APIs, such as character devices, file descriptors, etc. The functionality of this crate is mostly intended for use in LKMs but does provide some features that could be used elsewhere such as random numbers etc. The `alloc` and `core` crates constitute most of the `std` library

in Rust and respectively the implementation of a memory allocator and core functionality. The `alloc` and `core` crates are often used in embedded systems and other situations where no operating system to provide the functionality of the standard library. The `core` crate exposes basic functionality such as primitive types, references, etc. The `alloc` crate exposes memory allocations and in userspace uses some exposure of `malloc`, while in kernel space may use either `kmalloc` or `kvmalloc` to allocate physical and virtual memory inside the kernel. In its current form, the `alloc` crate does not provide much functionality. Only simple allocation types such as `Box` are exposed and their API is conservative. The reason behind this is that the kernel has no way to handle Out-Of-Memory cases. Thus most data structures are simply not allowed, because they do not expose a fallible way to allocate memory. Whenever a new allocation needs to happen a `try_new()` function can be called, which will return a `Result` type with either a reference or an error. For infallible memory allocations with `new()` an out-of-memory will throw an exception, which there is no good way to handle. The only data structure available is `Vec`, a dynamic array. For faster performance on lookup, we might need other data structures. Furthermore, the `alloc` crate is compiled with a `no_rc` feature meaning there is no way to use the reference counted pointers defined in Rust. The reason for this is that maintainers of the Rust functionality in Linux have decided that it is unnecessary, since the C part of the kernel already defines a reference counting functionality. To the best of my knowledge, there is no clear exposure of this functionality in any of the crates available. We need reference counting for our implementation. It is easy to remove this restriction but may make a potential PCC implementation harder to get merged into the upstream Linux.

It is possible to compile other crates than the ones defined above. The requirement is that the package must support a `no_std` feature, meaning it relies on `alloc` and `core` instead of `std`, and that also has no infallible memory allocations. One example of a library that does this is parser-combinator library `nom`, which we use for parsing.

2.3.2 Reference Counting

In the implementation, we make heavy use of reference counting, and we must therefore forego the `no_rc` restriction. We use compile time references when possible to not unnecessarily create new objects. When compile-time references are not possible, because we don't know the owner of a value and thus also not the lifetime of it, we instead use reference-counted pointers. Most of the functions we describe return values. The ownership will then lie at the caller of the function, but in some cases, the owner of a result value may be the context we do type-checking with respect to. This for instance happens when inferring the type of a variable. Furthermore, because of the lifetime guarantee, there is no way to create a value and return a reference to it.

The reference-counted smart pointer looks as follows:

```
pub struct Rc<T: ?Sized> {
    ptr: NonNull<RcBox<T>>,
    phantom: PhantomData<RcBox<T>>,
}
```

An `Rc` is nothing more than a struct that contains a pointer to the inner value that is referenced and a phantom field. The phantom field is merely there to keep strong static typing similar to a phantom type in Haskell. The `ptr` in this struct points to the following struct:

```
#[repr(C)]
struct RcBox<T: ?Sized> {
    strong: Cell<usize>,
    weak: Cell<usize>,
    value: T,
}
```

This contains the values and the counts for strong and weak reference counts. Whenever the `Rc` is cloned we simply take the `RcBox` inside of `Rc`, increment the pointer, and construct a new `Rc` struct. The ease of use then comes from the `Drop` trait which will either decrement the count in the `RcBox` and drop the `Rc` or it will drop both if the strong count is 0. Hence we can easily create new references without knowing the owner, as it does not matter since they are deallocated automatically. This gives some overhead compared to regular references but is highly likely more efficient than cloning.

2.4 Curry Howard Isomorphism

The Curry-Howard isomorphism states that a proposition is the type of its proofs, and is an equivalence between proof calculi and formal type systems. Concretely the isomorphism is a 2-level equivalence between logical formulas and types on one level, and logical proofs and computational programs on the other. For formulas and types, we have the following correspondence:

logic	type
\top (true)	()
\perp (false)	void
\wedge (conjunction)	product type
\vee (disjunction)	sum type
\rightarrow (implication)	function type
\forall (universal quantification)	Π (dependent type)
\exists (existential quantification)	Σ (dependent sum type)

Take for example the formula $A \wedge P \rightarrow P$. Informally, when we have evidence for both A and P then we can get a P , and similarly for types if we can give a pair (A, P) then we can get the inhabitant P .

For proof to such a proposition we consider the correspondence between natural deduction and lambda calculus, where

deduction rule	lambda term
hypothesis	free variable
implication elimination	application
implication introduction	abstraction

We then further consider destructors for sum and product types. For proposition such as $P \wedge (Q \vee R) \rightarrow ((P \wedge Q) \vee (P \wedge R))$: we can construct a program, in a Haskell-like syntax:

```
f :: (P, Either Q R) -> Either (P, Q) (P, R)
f (p, Left q) = Left (p, q)
f (p, Right r) = Right (p, r)
```

Given a pair of type $(P, \text{Either } Q \ R)$ we may construct a type $\text{Either } (P, Q) \ (P, R)$, hence constituting a proof for the proposition. One thing to note however is that we require programs to be total, as we could otherwise give evidence for the uninhabited void type. We will later see how we can construct more rigorous proofs using Π .

3 Proof Carrying Code in the Kernel

Before we can get a grasp of what it takes to make PCC work in the kernel, one must first have a better understanding of what is actually happening in the verifier. In Section 3.1 we take a “deep dive” into the inner workings of the verifier. The purpose of this is to create a bridge between eBPF and PCC and to further emphasize why we should investigate other methods of verification than the static analysis that exist at the moment. In Section 3.2 we describe an overall design of how we can use PCC in the Linux kernel. In this section, we further argue for some general design decisions taken in the process of creating the proof-checking part of the PCC system.

3.1 Spelunking the eBPF verifier

In this section, we describe in detail how an eBPF program is loaded. We first describe, how the `bpf` syscall is defined and we then proceed to give a general understanding of what steps are taken in the process of loading an eBPF program. Specifically, it is important what requirements the programs must follow and how this is realized.

3.1.1 The `bpf` syscall

All interaction between user and kernel space regarding eBPF-related matter uses the `bpf` syscall¹ and has the following signature:

```
asmlinkage long
sys_bpf(int cmd, union bpf_attr *attr, unsigned int size);
```

Argument `cmd` is an integer that defines the intended interaction. For the purpose of this project, we only care about the `cmd` `BPF_PROG_LOAD`, but intentions such as `BPF_MAP_UPDATE_ELEM`, `BPF_MAP_CREATE` and `BPF_MAP_LOOKUP_ELEM` also exists.

To be able to load an eBPF program one of the following criteria must be met:

- Either the caller of the syscall must be root or be bpf capable,
- or the `kernel.unprivileged_bpf_disabled` kernel parameter must be set to 0, meaning regular users are capable of loading programs. As mentioned previously for security reasons this is disabled in most major Linux distributions.

The `attr` argument is a union of structures that must correspond to the argument type. For program loading, `attr` notably contains the type of program to load, which could be socket

¹`bpf()` has syscall number 321

programs, kernel probes, Express Data Path, or one of the many other possibilities, as well as the program that needs to be loaded. The syscall will call the appropriate `cmd` after some sanity checks, such as the well-formedness of the `bpf_attr` union w.r.t the command and the `size` parameter. The first function called is `bpf_prog_load`.

3.1.2 Capabilities and Kernel configurations: `bfp_prog_load`

The main purpose of `bfp_prog_load` is to check for user capabilities and setting the parameters used in the verifier based on kernel configurations. It checks the following:

1. Normal users may only load socket programs, whilst network-related eBPF programs such as XDP requires network or system administrator capabilities. For performance monitoring the user loading the program must have `perfmon` capabilities.
2. Only `bpf_capable` users may use unaligned memory access in eBPF maps.
3. eBPF programs must be between 1 and 4096 instructions. For capable users, the limit is 1 million instructions.
4. The license of the program is checked and programs with GPL capabilities may call helper functions with GPL license.
5. Programs may be either eBPF programs or BTF objects, which is type-level information about eBPF programs. These are mainly used as debug information about eBPF programs and they are irrelevant to the verifier, thus also the work done in this report.
6. The program is checked for device boundness, as eBPF programs can be offloaded to devices.

These are all properties of an eBPF program, and must still be checked even if the verifier is not present. In the design of a new loading procedure `bpf_prog_load` should be mostly kept intact. We would however still need to modify the call to `bpf_check` as this is the entry point to the verifier.

3.1.3 Static analysis: `bpf_check`

The `bpf_check` is what we usually denote as the verifier. Firstly the checking environment is set up. The environment is a big struct with all the necessary information to complete the validation. The procedure starts by checking more capabilities, for instance, we have the following lines:

```
env->allow_ptr_leaks = bpf_allow_ptr_leaks();
env->allow_uninit_stack = bpf_allow_uninit_stack();
env->bypass_spec_v1 = bpf_bypass_spec_v1();
env->bypass_spec_v4 = bpf_bypass_spec_v4();
env->bpf_capable = bpf_capable();
```

The first 4 are flags only set for `perfmon` capabilities, which allow `perfmon` capable users to do more with the eBPF stack. Environment variables such as these introduce a dilemma in terms of what can be removed from the verifier. This is discussed in 3.2.2.

After these initial flags, the function does the following checks:

1. Firstly eBPF subprograms and kernel helper functions are added to the environment. Notice here, that the main eBPF program is also considered a subprogram, so whenever we state subprograms are checked this corresponds to the main as the eBPF subprograms.
2. Then function `check_subprogs` is called, where some simple checks are conducted, such as subprograms not being allowed to jump outside of their own address space. The last instruction of a subprogram must either exit or jump to a subprogram, constituting a tail call.
3. If the eBPF program is supposed to be device bound, it is prepared for that. We omit the details as we have not taken the device offloaded into account for this project. But it is an additional factor to consider in the replacement of the verifier.
4. Next `bpf_check` will check the control flow graph for loops using a non-recursive depth-first search approach. If a cycle is detected, the program is rejected.
5. All the subprograms are then checked according to their BPF TypeFormat (BTF).
6. The last step before loading a program is abstract interpretation using tri-state numbers (`tnum`).

The following is a simplification of the kernel documentation about the verifier[8]. A program must follow these requirements:

1. Registers may not be read unless they have previously been written. This is to ensure no kernel memory can be leaked.
2. Registers can either be scalars or pointers. After calls to kernel functions or when a subprogram ends, registers `r1-r5` are forgotten and thus cannot be read before written. `r6-r9` are callee saved and thus still available.
3. Reading and writing may only be done by registers marked by `ptr_to_ctx`, `ptr_to_stack`, or `ptr_to_map`. These are bound and alignment checked.
4. Stack space, for the same reason as registers, may not be read before it has been written.
5. External calls are checked at entry to make sure the registers are appropriate w.r.t. the external functions.
6. All read and writes to the stack and maps should be within bounds.
7. Division by 0 is not allowed unless the divisor is a register in which case the program is patched later in the verification process.

To keep track of this the verifier will do abstract interpretation. The verification process tracks minimum and maximum values in both the signed and unsigned domain. It furthermore uses `tnums` which is a pair of a mask and a value. The mask tracks unknown bits. Set bits in the value are known to be 1. The program is then traversed and updated modulo the instructions. For instance if register `r2` is a scalar and known to be in the range between $(0, \text{IMAX})$ then after abstractly interpreting a conditional jump `r2 > 42` the current state is split in two and the state where the condition is taken now have an updated range of $42 \leq r2 \leq \text{IMAX}$. Pointers are handled similarly, however since pointer arithmetic is inherently dangerous, modifying a pointer is very limited in eBPF. Additionally, pointers may be interpreted as

different types of pointers and are checked wrt. the program type they occur in. For instance `BPF_MAP_TYPE_SOCKMAP` may only be used with socket-type programs.

After abstract interpretation, the stack depth is checked, meaning we simply check if the function calls can fit within the stack space allocated for the eBPF program.

Next dead code is eliminated. The argumentation in the comments for the implementation is questionable. Specifically, they mention that malicious code can have dead code too, which clearly is correct, but also completely irrelevant. Especially since they are turned into `JAL` instructions.

If all these requirements are met, then an eBPF program is loaded. This mapping is simplified a lot, but it shows that the current process of checking a valid eBPF program has many steps of which some are directly code specific and some are tied to intentions and capabilities. This makes a PCC system a little more difficult to realize.

3.2 eBPF and PCC

From the description of PCC in 2.2 and the description of the eBPF subsystem above, it should be clear that eBPF and the verifier have some clear similarities with PCC. eBPF presents a set of security policies, which must be followed to load programs into the kernel. Likewise, there is a clear distinction between the user space and kernel space or in PCC terminology between code producer and consumer. Where the current eBPF loading system differs from the PCC described by Necula is where the responsibility lies. Considering the pipeline we have:

1. **Compilation and Certification:** For PCC the untrusted program is both compiled and a certificate for safety policy compliance is generated by the code producer. eBPF does not really “do” anything at this stage as source code is passed directly to the kernel using the syscall, and then possibly JIT compiled later in the pipeline.
2. **Verification of certificate:** In PCC the consumer will check the validity of the certification wrt. the safety policy and compatibility between the code and the certificate. eBPF will have to do a similar check but directly on the eBPF program. From a purely complexity-wise standpoint verification by checking a proof should not be any harder than performing the static analysis done by the verifier.
3. **Running:** In both structures, once a certificate is checked the program is free to use possibly many times.

So if they are so similar in structure, why would we want to replace the verifier with an actual proof checker? As already mentioned the eBPF verifier has been prone to bugs in the past, and the code for the verifier is characterized by patching of these bugs, instead of implementing a proven sound abstract interpretation. Having a proof checker that implements a sound approach to proof checking will give a higher certainty in safety as well as more accepted programs because a proof checker does not have to be as conservative as static analysis.

3.2.1 How to add PCC to eBPF

To realize PCC in the Linux kernel we must extend it in some way. There are mainly four ways we can extend the Linux kernel by the documentation[1].

1. If an operation can be achieved using one of the many filesystems already present in the kernel, then it should do so.
2. For device-specific operations, we should consider an LKM.
3. If new functionality is to the kernel it should be in the form of a syscall.
4. If strictly necessary a system call should be modified, but its API should be the same.

The first two options are not really viable solutions for a PCC infrastructure, since it would require plugging into the subsystem in some way, and this interaction in general seems to break many responsibility patterns and would still require modifying the `bpf` syscall. Option 3 likewise imposes a responsibility mismatch. All interactions with the eBPF subsystem go through the syscall, and having separate system calls to validate the code and load it is not optimal. Likewise, the proof checker cannot substitute the entire loading process but only the verifier, and probably not all of the checks can be completely removed. For instance, we might still require capabilities to be checked separately but include memory alignment in the verification condition for the program. On the other hand, capabilities are just boolean flags, which can easily be represented in logic. Such design decisions would require more experimentation.

The most optimal solution would be to modify the `bpf` syscall directly. We can here also get a partial transition by adding the feature of a proof checker and giving the proof as part of the `attr` struct for the syscall. Then when confidence in the proof checker is high enough the verifier can be phased out.

3.2.2 Certifications

As previously mentioned we only implement part of PCC in this report, and this implementation should also just be seen as a prototype.

For the general architecture, we consider a certification format based on formal logic of verification conditions based on predicate transformer semantics[6]. This is particularly useful for the automatic generation of formulas to describe programs. The process amount to showing the satisfiability of the negated formula, since this gives validity. The process of proving the validity of such a verification condition however is not a simple task and is, depending on the logic, undecidable. Checking the satisfiability of a formula can be done by a Satisfiability Modulo Theories (SMT) solver. In some SMT solvers, it is possible to extract a proof that a certain formula is satisfiable. We have in this work considered two output formats/languages, Alethe[alethe] and Logical Framework with Side Conditions (LFSC)[17], supported by the CVC5 SMT solver. I will briefly describe why I have chosen to use the LFSC language over Alethe.

Both formats are based on S-expressions and therefore simple to parse.

Alethe is designed to be easily readable by humans and structured as a box-style proof, with subproofs, conclusions, etc. which make understanding the proofs easy. This is not a property necessary for a PCC architecture where we want to automate the entire process and do not care about the plain text format, but in fact, rather would want a binary format. By this construction, the Alethe format provides a set of 91 inference rules [alethe], on which proofs are built. As an example, rule 20 is reflection often denoted as `refl` which states equality after applying a context. This entails that an implementation must implement all rules necessary for a security policy and will not be as easily extended in the future as part of a Linux kernel.

LFSC on the other hand is a metaframework that exploits the Curry-Howard isomorphism by dependent type theory. This meta-framework allows for the security policy to be established by signatures, which encodes similar rules to the ones defined in Alethe. The meta-framework allows us to implement a simple algorithm for type-checking signatures and verification conditions, which once defined does not have to be changed. New functionality can then easily be added by new signatures. Furthermore, this approach can move bugs out of the in-kernel certifier and into the specification. This will enable system administrators to quickly deploy fixes for a bug, by not allowing specific faulty signatures.

Another important feature of a certification checker in the kernel is that it should be both memory and time efficient. It is hard to consider both time and memory use of the two languages without actually implementing both of them. There already exists an implementation for each. LFSC has the *lfsc* proof checker distributed by CVC5[11], and Alethe has an unofficial proof checker called *caracara*[4]. We could use these for comparison, but *caracara* does not support proofs using bit-vectors, which is a must for a proof checker that needs to validate programs with bounded values. Hence we cannot say anything concrete about the performance of these compared to each other at the given time.

Overall LFSC provides a better basis for the use case. When using LFSC in the proof checker we suggest that the process of checking that the eBPF program to be loaded corresponds to the proof is done by generating a verification condition and then comparing it to the assertions in the proof. The specific formula that needs to be proved is directly embedded in an LFSC proof. Hence we suggest that an in-kernel verification condition generator generates formulas that can easily be compared with the formula from the proof. This also means that we can reuse some of the code used to implement the proof checker for the VC generation.

4 Logical Framework with Side Conditions

LFSC[17] is an extension of the Edinburgh Logical Framework (LF)[9] and is a predicative typed lambda calculus with dependent types. This allows proof systems to be encoded as signatures, which amounts to a set of typing declarations. LFSC extends LF by including side conditions. In this context side conditions refers to a functional programming language with an operational semantic which is evaluated during typechecking. Section 4.1 gives an introduction to the abstract syntax of LFSC. Section 4.2 describes the *signatures* and *contexts* under which typing is judged and then we describe the typing rules and the operational semantics of the side conditions in 4.3 and 4.4. Lastly we briefly introduce the concrete syntax in 4.5 and then we present a small example that shows that $P \wedge \neg P$ is unsatisfiable using LFSC.

4.1 Abstract Syntax

Figure 3 shows the 5 categories of the LFSC language. At its core LFSC is a typed lambda calculus and it consists of terms/objects, types, and kinds. The last two categories are patterns and side conditions.

Terms are denoted by M , N , and O and are used as syntactical entities, proofs, or inference rules in a logic. Types are denoted by A and B and are used for the classification of terms and to describe judgments and assertions. Kinds are denoted by K and are used to classify types.

$$\begin{aligned}
K &::= \mathbf{type} \mid \mathbf{type}^c \mid \mathbf{kind} \mid \Pi x : A.K \mid \mathbf{int} \mid \mathbf{rational} \\
A, B &::= a \mid A M \mid \Pi x : \{S M\}.A \mid \Pi x : A.B \\
M, N, O &::= x \mid c \mid z \mid q \mid * \mid M : A \mid \mathbf{let} \ x \ M \ N \mid \lambda x.M \mid \lambda x : A.M \mid M \ N \\
P &::= c \mid c \ x_1 \dots x_n \\
S, T, U &::= x \mid c \mid - S \mid S \oplus S \mid c \ S_1 \dots S_n \mid \mathbf{let} \ x \ S \ T \mid \mathbf{markvar} \ S \mid \\
&\quad \mathbf{ifequal} \ S_1 \ S_2 \ T \ U \mid \mathbf{match} \ S \ (P_1 \ T_1) \dots (P_n \ T_n) \mid \mathbf{fail} \ S \mid \\
&\quad \mathbf{ifneg} \ S \ T \ U \mid \mathbf{ifzero} \ S \ T \ U \mid \mathbf{ifmarked} \ S \ T \ U \mid \mathbf{ztoq} \ S \\
\oplus &\in \{+, /, *\}
\end{aligned}$$

Figure 3: Syntactical categories of LFSC

We use x to be a metavariable ranging over the set of variables that might occur in terms, c to denote constants in terms, i.e. free variables. a will range over constants in types. Side conditions are denoted by S and T and U . Patterns describe patterns in side condition match cases and are denoted by P . Primitives and keywords of the side condition language are represented in **bold**. z and q are meta symbols representing integers and rationals. $*$ represents a hole, which can be filled with any term in the type-checking process. Both Π and λ are abstractions that bind the variable in the body. type annotations are given by $:$ and $\{S M\}$ is a pair.

4.2 Signatures and Contexts

In LFSC there are two constructs we use to keep track of variables and constants. We have signatures and contexts. Signatures are used to assign kinds and types to constants, thus defining the formal system on which terms are judged. Contexts are used to assign types to variables. we write them as in Figure 4 and use we use Σ , Σ' to denote the concatenation of the two signatures Σ and Σ' and similarly for contexts.

$$\begin{aligned}
\Sigma &::= \langle \rangle \mid \Sigma, a : K \mid \Sigma, c : A \\
\Gamma &::= \langle \rangle \mid \Gamma, x : A
\end{aligned}$$

Figure 4: Signatures and Contexts

The type system of LFSC is syntax-directed meaning there exists only a single typing rule for each syntactical object. We achieve this by bidirectional typing. That means instead of stating that an expression must have a type, we can either construct a type from it (called inference)

or we can check that an expression has a type. All assertions have one of the following forms.

$$\begin{array}{ll}
\Sigma \checkmark & (\Sigma \text{ is a valid signature}) \\
\vdash_{\Sigma} \Gamma & (\Gamma \text{ is a valid context in } \Sigma) \\
\Gamma \vdash_{\Sigma} K & (K \text{ is a kind in } \Gamma \text{ and } \Sigma) \\
\Gamma \vdash_{\Sigma} M \Leftarrow A & (M \text{ can be checked to have type } A \text{ in } \Gamma \text{ and } \Sigma) \\
\Gamma \vdash_{\Sigma} M \Rightarrow A & (M \text{ can be inferred to have type } A \text{ in } \Gamma \text{ and } \Sigma)
\end{array}$$

The notion of inference here describes the type-level semantic of an object, meaning if there is a judgment from M that gives A , then $M \Rightarrow A$. The notion of checking a type we consider to be $\beta\eta$ equivalence. Specificially $(M \Leftarrow A) \equiv (M : B \wedge B \equiv_{\beta\eta} A)$.

The validity of signatures is depicted in Figure 5. The empty signature is valid. The concatenation of signatures is valid if Σ is valid and the constant is not present in the context. For *KIND-SIG* K must be valid in Σ , while for *TYPE-SIG* we require that A can be inferred to have type K .

$$\begin{array}{c}
\text{EMPTY-SIG} \frac{}{\langle \rangle \checkmark} \quad \text{KIND-SIG} \frac{\Sigma \checkmark \quad \vdash_{\Sigma} K \quad a \notin \text{dom}(\Sigma)}{\Sigma, a : K \checkmark} \\
\text{TYPE-SIG} \frac{\Sigma \checkmark \quad \vdash_{\Sigma} A \Rightarrow K \quad c \notin \text{dom}(\Sigma)}{\Sigma, c : A \checkmark}
\end{array}$$

Figure 5: Valid signatures

Figure 6 depicts the validity of contexts. An empty context is valid under Σ given the validity of Σ . For concatenation, x must not occur in Σ but can occur in Γ , and again we require that A can be inferred as any kind.

$$\begin{array}{c}
\text{EMPTY-CTX} \frac{\Sigma \checkmark}{\vdash_{\Sigma} \langle \rangle} \quad \text{TYPE-CTX} \frac{\vdash_{\Sigma} \Gamma \quad \Gamma \vdash_{\Sigma} A \Rightarrow K \quad x \notin \text{dom}(\Sigma)}{\vdash_{\Sigma} \Gamma, x : A}
\end{array}$$

Figure 6: Valid contexts

4.3 Typing

With the signature and contexts we can define the typing rules for LFSC. For the entire section we brevitate \vdash_{Σ} to \vdash and it is assumed Σ is valid.

4.3.1 Lookup

Figure 7 describes type inference for variables, constants and built in types. That is, they contain all rules where variables or constants are inferred w.r.t. the signature and context. Given that Γ and Σ is valid, we can infer the built in types **type** and **type^c** to be **kind**. Notice here that **type** and **type^c** are kinds which describes types, whereas **kind** describes kinds. Notice further that **kind** cannot be inferred and thus ensures that no type contains itself hence providing a consistent metalogic. Constants can be inferred either as a kind or a type respectively if they appear in the signature, whereas object level variables must occur in the context.

$$\begin{array}{c}
\text{TYPE} \frac{\vdash \Gamma}{\Gamma \vdash \mathbf{type} \Rightarrow \mathbf{kind}} \quad \text{TYPEc} \frac{\vdash \Gamma}{\Gamma \vdash \mathbf{type}^c \Rightarrow \mathbf{kind}} \\
\text{MPZ} \frac{\vdash \Gamma}{\Gamma \vdash \mathbf{int} \Rightarrow \mathbf{type}} \quad \text{MPQ} \frac{\vdash \Gamma}{\Gamma \vdash \mathbf{rational} \Rightarrow \mathbf{type}} \\
\text{LOOKUP-CTX} \frac{\vdash \Gamma \quad x : A \in \Gamma}{\Gamma \vdash x \Rightarrow A} \\
\text{LOOKUP-KIND-SIG} \frac{\vdash \Gamma \quad a : K \in \Sigma}{\Gamma \vdash a \Rightarrow K} \quad \text{LOOKUP-TYPE-SIG} \frac{\vdash \Gamma \quad c : A \in \Sigma}{\Gamma \vdash c \Rightarrow A}
\end{array}$$

Figure 7: Typing rules for looking up types.

4.3.2 Terms and Types

Figure 8 describes the bidirectional typing of LFSC.

$$\begin{array}{c}
\text{PI} \frac{\Gamma \vdash A \Leftarrow \mathbf{type} \quad \Gamma, x : A \vdash C \Rightarrow \alpha \quad \alpha \in \{\mathbf{type}, \mathbf{type}^c, \mathbf{kind}\}}{\Gamma \vdash \Pi x : A. C \Rightarrow \alpha} \\
\text{PI-SC} \frac{\Gamma \vdash S \Rightarrow \mathbf{type} \quad M \Rightarrow \mathbf{type} \quad \Gamma \vdash B \Rightarrow \mathbf{type}}{\Gamma \vdash \Pi x : \{S \ M\}. B \Rightarrow \mathbf{type}^c} \\
\text{TYPE-APP} \frac{\Gamma \vdash A \Rightarrow \Pi x : B. K \quad \Gamma \vdash M \Leftarrow B}{\Gamma \vdash A \ M \Rightarrow [M/x]K} \quad \text{APP-HOLE} \frac{\Gamma \vdash M \Rightarrow \Pi x : A. B}{\Gamma \vdash M * \Rightarrow [* / x]B} \\
\text{APP} \frac{\Gamma \vdash M \Rightarrow \Pi x : A. B \quad \Gamma \vdash N \Leftarrow A}{\Gamma \vdash M \ N \Rightarrow [N/x]B} \quad \text{ANN} \frac{\Gamma \vdash M \Leftarrow A}{\Gamma \vdash M : A \Rightarrow A} \\
\text{APP-SC} \frac{\Gamma \vdash M \Rightarrow \Pi x_1 : A. (\Pi x_2 : \{S \ O\}. B) \quad \Gamma \vdash N \Leftarrow A \quad |\Sigma| \vdash \epsilon; [N/x]S \downarrow [N/x]O; \sigma}{\Gamma \vdash M \ N \Rightarrow [N/x_1]B} \\
\text{LAM-ANN} \frac{\Gamma \vdash A \Rightarrow \mathbf{type} \quad \Gamma, x : A \vdash M \Rightarrow B}{\Gamma \vdash \lambda x : A. M \Rightarrow \Pi x : A. B} \quad \text{Q} \frac{\vdash \Gamma}{\Gamma q \Rightarrow \mathbf{mpq}} \\
\text{LAM} \frac{\Gamma, x : A \vdash M \Rightarrow B}{\Gamma \vdash \lambda x. M \Leftarrow \Pi x : A. B} \quad \text{Z} \frac{\vdash \Gamma}{\Gamma \vdash z \Rightarrow \mathbf{mpz}}
\end{array}$$

Figure 8: Bidirectional typing rules for LFSC

Here $[M/x]K$, denotes the substitution of M with x in K . The letter C is either a type or a kind. $|\Sigma|$ denotes all side condition function definitions in Σ .

- *ANN* states that given object M can be checked to have type A , then the annotation can be inferred to type A .
- *LAM-ANN* states that a bound variable x has type A in M can be inferred as a $\Pi x : A. B$, given that A is **type** and that M can be inferred as B with $x : A$ added to the environment.
- Lambda abstractions, *LAM*, is the only type that we cannot directly infer, but instead must be checked to be a function type. A lambda is a valid function type if the body M can be inferred as B with $x : A$ added to the context.
- Both *TYPE-APP* and *APP* states that if the function point (left construct) is a function type and that the argument (right construct) can be checked to be the type of domain of

the function type, then an application can be inferred as the substitution of the operand with the x in the construct describing the range of the function.

- *APP-HOLE* is the only rule where a hole might occur. It may only occur as argument in an application. We don't type check it as it is trivially the correct term, but the first time a hole occurs as part of a check it will be filled with the other value, continuing to have the type of the other term.
- *APP-SC* is where the type system of LFSC differs from that of LF. If the type of M is a function type where the range itself is a function type with a side condition as domain, then N will be checked to have type A and the side condition S is evaluated by its operational semantics, and must result in O .

4.3.3 Side conditions

For easier readability we split the side condition language into three subcategories:

1. **numerical** side conditions that focused around numerical values.
2. **side effects** side conditions that can update the state in evaluation.
3. **compound** side conditions that are construct that don't fall into the other categories.

For **numerical sideconditions** the rules are straight forward. They are represented for **integer** in Figure 9 but work similarly for **rational**, except for Z-TO-Q, which as the name suggest requires S to be of type **integer**. *IFNEG* and *IFZERO* are branching constructs where the condition must be one of the two number types and the branches must have the same type.

$$\begin{array}{c}
\text{BINOP} \frac{\Gamma \vdash S \Rightarrow \mathbf{int} \quad T \Rightarrow \mathbf{int}}{\Gamma \vdash S \oplus T \Rightarrow \mathbf{int}} \oplus \in \{+, *, /\} \quad \text{Z-TO-Q} \frac{\Gamma \vdash S \Rightarrow \mathbf{int}}{\Gamma \vdash \mathbf{ztoq} S \Rightarrow \mathbf{rational}} \\
\\
\text{IFNEG} \frac{\Gamma \vdash S \Rightarrow \mathbf{int} \quad \Gamma \vdash T \Rightarrow A \quad \Gamma \vdash U \Rightarrow A}{\Gamma \vdash \mathbf{ifneg} S T U \Rightarrow A} \quad \text{NEG} \frac{\Gamma \vdash S \Rightarrow \mathbf{int}}{\Gamma \vdash -S \Rightarrow \mathbf{int}} \\
\\
\text{IFZERO} \frac{\Gamma \vdash S \Rightarrow \mathbf{int} \quad \Gamma \vdash T \Rightarrow A \quad \Gamma \vdash U \Rightarrow A}{\Gamma \vdash \mathbf{ifzero} S T U \Rightarrow A} \quad \text{INT} \frac{}{\Gamma \vdash n \Rightarrow \mathbf{int}}
\end{array}$$

Figure 9: Typing rules for numerical sideconditions

For **side effects** rules we have that *LET-SC* works as its counter part, and *DO* is mere syntactical sugar for *LET* but with the binding of a variable, as can be seen from Figure 10 *MARKVAR* will simply return the type of the inner side condition and *IFMARKED* will check the “marked” side condition S and check that the two branches have the same type, with the type of the branches as the resulting type.

$$\begin{array}{c}
\text{DO} \frac{\Gamma \vdash S \Rightarrow A \quad \Gamma \vdash T \Rightarrow B}{\Gamma \vdash \mathbf{do} S T \Rightarrow B} \quad \text{IFMARKED} \frac{\Gamma \vdash S \Rightarrow A \quad T \Rightarrow B \quad U \Rightarrow B}{\Gamma \vdash \mathbf{ifmarked} n S T U \Rightarrow B} \\
\\
\text{Let} \frac{\Gamma \vdash S \Rightarrow A \quad \Gamma, x : A \vdash T \Rightarrow B}{\Gamma \vdash \mathbf{let} x S T \Rightarrow B} \quad \text{MARKVAR} \frac{\Gamma \vdash S \Rightarrow A}{\Gamma \vdash \mathbf{markvar} n S \Rightarrow A}
\end{array}$$

Figure 10: Typing rules for side effects

compound sideconditions is seen in Figure 11 may be a **fail**, described by the *FAIL* rule, which simply typechecks the inner value. The reason **fail** must take an argument is that we have no polymorphic type $\forall \alpha. \alpha \rightarrow \alpha$ under dependent types. For *IFEQ* S_1 and S_2 must have the same type and the branches T_1 and T_2 must equally have the same type. Match statements work similar to other functional languages; given the scrutinee S can be inferred as A then for all match cases P_i must also be inferreable as A , whilst all branches T_i must be inferreable to the same type B . $\text{ctx}(P_i)$ describes the context created from P_i . Concretely $\text{ctx}(P_i) = \langle \rangle, x_1 : A_1, x_2 : A_2, \dots, x_n : A_n$, when $P_i = cx_1 \ x_2 \ \dots \ x_n$. For applications the function point must not be a dependent function.

$$\begin{array}{c}
\text{SCAPP} \frac{\Gamma \vdash S \Rightarrow \Pi x : A.B \quad \Gamma \vdash T \Rightarrow C \quad x \notin FV(B)}{\Gamma \vdash S \ T \Rightarrow B} \quad \text{FAIL} \frac{\Gamma \vdash A \Rightarrow \text{type}}{\Gamma \vdash \text{fail } A \Rightarrow A} \\
\\
\text{MATCH} \frac{\Gamma \vdash S \Rightarrow A \quad \forall i \in \{1, \dots, n\}. (\Gamma \vdash P_i \Rightarrow A \quad \Gamma, \text{ctx}(P_i) \vdash T_i \Rightarrow B)}{\Gamma \vdash \text{match } S \ (P_1, T_1) \dots (P_n, T_n) \Rightarrow B} \\
\\
\text{IFEQ} \frac{\Gamma \vdash S_1 \Rightarrow A \quad \Gamma \vdash S_2 \Rightarrow A \quad \Gamma \vdash T_1 \Rightarrow B \quad \Gamma \vdash T_2 \Rightarrow B}{\Gamma \vdash \text{ifequal } S_1 \ S_2 \ T_1 \ T_2 \Rightarrow B}
\end{array}$$

Figure 11: Typing rules for compound side conditions

4.4 Operational Semantics of Side conditions

Now we present the operational semantics of the side condition language. It is shown in Figure 12. The operational semantics is under the judgment of $\Delta \vdash \sigma_1; S \downarrow T; \sigma_n$, where Δ describes all program definitions, and σ_1 and σ_n describes states mapping symbols to markings. Here σ_1 is the state S is evaluated under, and σ_n is the state where S has been evaluated to T . Markings are simply a collection of 32 boolean flags, which can then be used in **ifmarked** conditions. The Δ is elided in all cases where it is unused.

Errors are not included in the operational semantics. Errors might occur when **fail** is evaluated, a scrutinee does not match any pattern, a **markvar** or **ifmarked** does not evaluate to a variable or if division by 0 occurs.

For a brief rundown of the rules we have:

- *CST-O*, *VAR-O*, and *NUM-O* simply evaluates to themselves and the store is unchanged.
- *CST-APP* applies a constant to n side conditions, updates the store with respect to all of them, and the resulting value is the constant applied to each updated S' .
- *LET-O* and *DO-O* evaluate S , then evaluates T with the updated store and, in the case of *LET-O*, substitute occurrences of x in T .
- The two rules *IFEQUAL-T* and *IFEQUAL-T* describes a standard semantic for equality checks. Two terms S'_1 and S'_2 are considered equivalent with respect to $\beta\eta$ -equivalence.
- Match constructs evaluate the scrutinee S and match the result with one of the patterns. If a pattern matches then the given branch is evaluated.
- *FUN-APP* refers to the application of a side condition program. f must be a program in Δ . All arguments are evaluated and then the results are substituted into the body of the program T which is finally evaluated.

- *BINOP-O* and *NEG-O* work similarly to any other language.
- *ZTOQ-O* evaluate S to an integer z and then make the rational z/z .
- *IFNEG* and *IFZERO* rules are very similar; based on the evaluation of S , either branch T or branch U is evaluated.
- *IFMARKED* is again similar however the branching depends on the marking of variable x in the store.
- *MARKVAR-O* is the only rule that can update the store, by simply switching the flag for the specific mark.

$$\begin{array}{c}
\text{CST-APP} \frac{\forall i \in \{1, \dots, n\}. (\sigma_i; S_i \downarrow S'_i; \sigma_{i+1})}{\sigma_1; (c \ S_1 \dots S_n) \downarrow (c \ S'_1 \dots S'_n); \sigma_{n+1}} \\
\text{LET-O} \frac{\sigma_1; S \downarrow S'; \sigma_2 \quad \sigma_2; [S'/x]T \downarrow T'; \sigma_3}{\sigma_1; (\text{let } x \ S \ T) \downarrow T'; \sigma_3} \quad \text{DO-O} \frac{\sigma_1; S \downarrow S'; \sigma_2 \quad \sigma_2; T \downarrow T'; \sigma_3}{\sigma_1; (\text{do } S \ T) \downarrow T'; \sigma_3} \\
\text{IFEQUAL-T} \frac{\sigma_1; S_1 \downarrow S'_1; \sigma_2 \quad \sigma_2; S_2 \downarrow S'_2; \sigma_3 \quad S'_1 \equiv S'_2 \quad \sigma_3; T_1 \downarrow T'_1; \sigma_4}{\sigma_1; (\text{ifequal } S_1 \ S_2 \ T_1 \ T_2) \downarrow T'_1; \sigma_4} \\
\text{IFEQUAL-F} \frac{\sigma_1; S_1 \downarrow S'_1; \sigma_2 \quad \sigma_2; S_2 \downarrow S'_2; \sigma_3 \quad S'_1 \not\equiv S'_2 \quad \sigma_3; T_1 \downarrow T'_1; \sigma_4}{\sigma_1; (\text{ifequal } S_1 \ S_2 \ T_1 \ T_2) \downarrow T'_1; \sigma_4} \\
\text{MATCH-O} \frac{\sigma_1; S \downarrow (c \ S_1 \dots S_m); \sigma_2 \quad \exists i. P_i = (cx_1 \dots x_m) \quad \sigma_2; [S'_1/x_1, \dots, S'_m/x_m]T_i \downarrow T'; \sigma_3}{\sigma_1; (\text{match } S \ (P_1 T_1) \dots (P_n T_n)) \downarrow T'; \sigma_3} \\
\text{FUN-APP} \frac{\forall i \in \{1, \dots, n\}. (\Delta \vdash \sigma_i; S_i \downarrow S'_i; \sigma_{i+1}) \quad (f(x_1 : A_1 \dots x_n : A_n) : B = T) \in \Delta \quad \Delta \vdash \sigma_{n+1}; [S'_1/x_1, \dots, S'_n/x_n]T \downarrow T'; \sigma_{n+2}}{\sigma_1; (f \ S_1 \dots S_n) \downarrow T'; \sigma_{n+2}} \\
\text{BINOP-O} \frac{\sigma_1; S \downarrow r_1; \sigma_2 \quad \sigma_2; T \downarrow r_2; \sigma_3 \quad r = r_1 \oplus r_2}{\sigma_1; S \oplus T \downarrow r; \sigma_3} \oplus \in \{+, *, /\} \\
\text{IFNEG-T} \frac{\sigma_1; S \downarrow r; \sigma_2 \quad r < 0 \quad \sigma_2; T \downarrow T'; \sigma_3}{\sigma_1; (\text{ifneg } S \ T \ U) \downarrow T'; \sigma_3} \quad \text{CST-O} \frac{}{\sigma_1; c \downarrow c; \sigma_1} \\
\text{IFNEG-F} \frac{\sigma_1; S \downarrow r; \sigma_2 \quad r \geq 0 \quad \sigma_2; U \downarrow U'; \sigma_3}{\sigma_1; (\text{ifneg } S \ T \ U) \downarrow U'; \sigma_3} \quad \text{VAR-O} \frac{}{\sigma_1; x \downarrow x; \sigma_1} \\
\text{IFZERO-T} \frac{\sigma_1; S \downarrow r; \sigma_2 \quad r = 0 \quad \sigma_2; T \downarrow T'; \sigma_3}{\sigma_1; (\text{ifzero } S \ T \ U) \downarrow T'; \sigma_3} \quad \text{NUM-O} \frac{}{\sigma_1; r \downarrow r; \sigma_1} \\
\text{IFZERO-F} \frac{\sigma_1; S \downarrow r; \sigma_2 \quad r \neq 0 \quad \sigma_2; U \downarrow U'; \sigma_3}{\sigma_1; (\text{ifzero } S \ T \ U) \downarrow U'; \sigma_3} \quad \text{NEG-O} \frac{\sigma_1; S \downarrow r; \sigma_2}{\sigma_1; -S \downarrow r; \sigma_2} \\
\text{ZTOQ-O} \frac{\sigma_1; S \downarrow z; \sigma_2 \quad r = z/1}{\sigma_1; \text{ztoq } S \downarrow r; \sigma_2} \quad \text{MARKVAR-O} \frac{\sigma_1; S \downarrow x; \sigma_2}{\sigma_1; (\text{markvar } S) \downarrow x; \sigma_2[x \mapsto \neg \sigma_2 x]} \\
\text{IFMARKED-T} \frac{\sigma_1; S \downarrow x; \sigma_2 \quad \sigma_2 \ x \quad \sigma_2; T \downarrow T'; \sigma_3}{\sigma_1; (\text{ifmarked } S \ T \ U) \downarrow T'; \sigma_3} \\
\text{IFMARKED-F} \frac{\sigma_1; S \downarrow x; \sigma_2 \quad \neg(\sigma_2 \ x) \quad \sigma_2; U \downarrow U'; \sigma_3}{\sigma_1; (\text{ifmarked } S \ T \ U) \downarrow U'; \sigma_3}
\end{array}$$

Figure 12: Operational semantics for side conditions

4.5 Concrete Syntax

Although the concrete syntax of LFSC is not terribly important, there is no clear presentation of how signatures are supposed to represent a formal system purely from the abstract syntax and the typing rules, as there is no way to extend it from the rules. We give a brief introduction to the examples presented later easier to understand.

LFSC implements the core language and side conditions as S-expressions. At the toplevel LFSC allows the following commands.

- **define** takes two arguments: the constant to be bound c and a term M . It will then bind c to the term M with type A , where A is the type of inference.
- **declare** takes a constant a and a type A , check that A is a kind, then bind a to A .
- **function** is used to define sideconditions. It takes an identifier, a pairwise list of (x, A) , which is the arguments to the function, then a return type and a body. It will check the body with the parameters added to a context, and then match the type of the program body with the return type.
- **check** will take a single argument, a term or a type, and then typecheck it.

The intention here is that **declare** and **function** is allowed only in the signature definitions, which defines the formal system, and then a user can utilize **define** and **check** to construct a proof under the formal system defined in the signature.

For the term language the following symbols are used:

Abstract Syntax	Concrete Syntax
$\Pi x : A . B$	<code>! x A B</code>
$M : A$	<code>: A M</code>
$\lambda x : A . M$	<code># x A M</code>
$\lambda x . N$	<code>\ x N</code>
<code>let x M N</code>	<code>@ x M N</code>
$\{ S T \}$	<code>^ S T</code>
<code>*</code>	<code>-</code>

The side condition language directly uses the keywords marked with bold in the previous sections.

4.6 Using LFSC to show $P \wedge \neg P$ is unsatisfiable

In this section, we give an example of how LFSC can be used to construct proofs. We do so, by explaining a proof that $P \wedge \neg P$ is unsatisfiable. The proof can be encoded by the following commands:

The program first defines a variable `cvc.p` by application of the function `var`, a function used to define free constants under SMT-Lib, and it is specifically characterized by a unique number and a sort. This makes `cvc.p` unique. `cvc.p` corresponds to the P in the formula above.

```

(define cvc.p (var 0 Bool))
(check
 (# a0 (holds (and cvc.p (and (not cvc.p) true)))
 (: (holds false)
 (resolution _ _ _
 (and_elim _ _ 1 a0)
 (and_elim _ _ 0 a0) ff cvc.p))))

```

Figure 13: Unsatisfiability proof of $P \wedge \neg P$ in LFSC

The proof is constructed by a check command. We first define `a0` stating that $P \wedge (\neg P \wedge \top)$ holds. The reason `true` is also included is because SMT-lib considers applications as *n*-ary functions and these will be represented as a null-terminated curried form of higher-order application. Specifically `and` is defined by:

```

1 (declare apply (! t1 term (! t2 term term)))
2 (declare f_and term)
3 (define and (# t1 term (# t2 term (apply (apply f_and t1) t2))))

```

The body of the check is then an annotation stating that line 5-7 should have the type `holds false`, meaning we can derive \perp and unsatisfiable. We notice that there are quite several holes. These will get filled out as we go.

`resolution` and `and_elim` are declared as follows:

```

1 (declare resolution (! c1 term
2                     (! c2 term
3                     (! c term
4                     (! p1 (holds c1)
5                     (! p2 (holds c2)
6                     (! pol flag
7                     (! l term
8                     (! r (^ (sc_resolution c1 c2 pol l) c) (holds c))))))))))
9 (declare and_elim (! f1 term
10                  (! f2 term
11                  (! n mpz
12                  (! p (holds f1)
13                  (! r (^ (nary_extract f_and f1 n) f2) (holds f2))))))

```

We can from these declarations see that the 3 holes on line 5 in Figure 13, should match parameters `c1`, `c2`, and `c`. These occur later in the type and can therefore be derived from the latter arguments. Similarly for `and_elim` the hole used for argument `f1` will be filled by argument `p` because the types must match.

Since the fourth argument of the `and_elim` applications are `a0` we can syntactically compare `a0` and `(holds f1)`, letting `f1 = (and cvc.p (and (not cvc.p) true))`. Notice here that both applications of `and_elim` are provided with 4 arguments, even though the type suggests, that it should take 5. However, from the typing rules the inner side condition of Figure 8 will be evaluated, hence when all types are checked `nary_extract` is run. `nary_extract` is defined as:

```

1 (function nary_extract ((f term) (t term) (n mpz)) term
2   (match t
3     ((apply t1 t2)
4      (mp_ifzero n
5        (getarg f t1)
6        (nary_extract f t2 (mp_add n (mp_neg 1))))))
7 )

```

It will extract the n -th element of an f application in t . so when `nary_extract` is called with `f_and`, `f1 = (and cvc.p (and (not cvc.p) true))` and `1`, and we by beta reduction have: `f1 = (apply (apply f_and cvc.p) (and (not cvc.p) true))` the only branch of `nary_extract` matches this scrutinee and we check if n is `0`. Since this is false we recursively call `nary_extract`. In the next iteration, we get: `t = (apply (apply f_and (not cvc.p)) true)`, now n is also `0` and we call `getarg` defined as follows:

```
1 (function getarg ((f term) (t term)) term
2   (match t ((apply t1 t2) (ifequal t1 f t2 (fail term))))))
```

Now `t = (apply f_and (not cvc.p))`. Since `t1` and `f` is both `f_and` we therefore get `(not cvc.p)` back.

This is then checked against `f2` (in the run case of `and_elim`) and the body `(holds f2)` is then checked. By a similar approach we get `cvc.p` from the other application of `and_elim`.

By now we have established all the arguments to `resolution`, `c1 = (not cvc.p)`, `c2 = cvc.p`, `c` is still a hole and `p1` and `p2` is `(holds (not cvc.p))` and `(holds cvc.p)` respectively. Finally `sc_resolution` is evaluated.

```
1 (function sc_resolution
2   ((c1 term) (c2 term) (pol flag) (l term)) term
3   (nary_elim f_or
4     (nary_concat f_or
5       (nary_rm_first_or_self f_or
6         (nary_intro f_or c1 false)
7         (ifequal pol tt 1 (apply f_not 1))
8         false)
9       (nary_rm_first_or_self f_or
10        (nary_intro f_or c2 false)
11        (ifequal pol tt (apply f_not 1) 1)
12        false)
13      false)
14    false))
```

At the innermost applications, we have calls to `nary_intro` which lifts a value into n -ary form. The two calls will turn `c1` into `(or c1 false)` and vice versa. `nary_rm_first_or_self` will then check if the result of `nary_intro` is equivalent to the `ifequal` call and return the fourth argument `false` if that is the case otherwise it will remove the first occurrence of `f_or`. For the application in line 4 the arguments are `(or (not cvc.p) false)` `(apply f_not cvc.p)` `false`. The first two arguments are not equivalent and thus we remove the first occurrence of `f_or` in the first argument. β -reducing the term results in `(apply (apply f_or (not cvc.p)) false)`, hence the result will be `false`. For line 5 we equally get `false`.

`nary_concat` will also return `false`, since the arguments are not n -ary applications. We can then clearly not eliminate any `f_or` from the expression, thus the result becomes `false`.

We match hole `f_2` in `resolution` with the result and get `holds false` which is equivalent to the annotated type, hereby concluding the proof.

From this small example the takeaway should be that side conditions can be useful because they allow recursive structure, meaning we for instance can extract the n^{th} clause of a conjunction, instead of applying conjunction elimination n times.

5 The in-kernel proof checker - LFSCR

In this section, we present the implementation, *lfscr*. We first describe some high-level design decisions taken to reduce the amount of memory needed and to make the type-checking algorithm efficient. We do so in Section 5.1. In Section 5.2 we discuss how to read and format the proofs to a format appropriate for the Linux kernel. Lastly, we go through the implementation of the type checker in Section 5.3

5.1 General implementation choices

In this section, we describe some general considerations for the implementation w.r.t performance. In short, LFSCR can use normalization by evaluation with De Bruijn indices and explicit substitutions as techniques.

5.1.1 De Bruijn Indices

From the description of the type semantics in 4.3, we notice that for a type to be correctly checked it must have definitional equality. Using De Bruijn indices makes this process a lot easier since it allows for considering α -equivalence as syntactical equivalence. Specifically with De Bruijn indices, when $\beta\eta$ normal form has been achieved, we get α -equivalence for free. De Bruijn indices further makes the process of beta-reduction easier, as variables don't have to be renamed, i.e. α -conversion. Consider application $(\lambda x. \lambda y. xy)y$. If we were to do direct substitution in $[y/x](\lambda y. xy)$ we would get $\lambda y. yy$. This changes the meaning of the term. De Bruijn indices instead swap each bound variable with a positive integer. The meaning of the integer n is then constituted by the n^{th} enclosing abstraction, Π , λ or *let*. Consider $\lambda x. \lambda y. \lambda y. xy$ and $\lambda y. \lambda x. \lambda x. yx$ which are on $\beta\eta$ -long form and α -equivalent but not syntactically identical. Using De Bruijn notation we get: $\lambda\lambda\lambda 2 0$ for both, since the function point in the inner application is described by the second innermost binder, starting from zero, whilst the argument for the application is 0, since it is captured by the innermost abstraction. If we again consider $(\lambda x. \lambda y. xy)y$ the De bruijn representation is $(\lambda\lambda 1 0)y$ and by beta reduction becomes $\lambda y 0$. We only consider De Bruijn notation for bound variables, thus we get around any capture avoiding complications. We could potentially also consider using De Bruijn indices for free variables, however this would complicate the code as this would require lifting binders. Although it could be interesting to consider De Bruijn levels since these are not relative to the scope.

We also consider De Bruijn indices for other binders such as program definitions:

```
(function sc_arith_add_nary ((t1 term) (t2 term)) term
  (a.+ t1 t2))
```

This program definition will get converted into:

```
(function sc_arith_add_nary ((term) (term)) term
  (a.+ _1 _0))
```

Where `_` prefixes De Bruijn index to distinguish them from integers.

Similarly in pattern matching, when a constructor is applied to multiple arguments, the following

```
(match t
  ((apply t1 t2)
    (let t12 (getarg f t1)
      (ifequal t12 1 tt (nary_ctn f t2 1))))
  (default ff))
```

gets converted into:

```
(match t
  ((apply 2)
    (let (getarg f _1)
      (ifequal _0 1 tt (nary_ctn f _1 1))))
  (default ff))
```

Notice here that the arguments `t1` and `t2` are substituted by `2` as argument for the constructor as we need to save the number of arguments to not lose the meaning of the constructor, as it must be fully applied. In the example we also converted the binder of the “let”.

5.1.2 Explicit substitutions

We have already touched upon substitution, but another matter at which we shall consider is the sheer cost of direct substitution. Performing direct substitution on terms we can cause an explosion in the size of the term and unnecessarily waste both memory and execution time because we have to copy the struct at each occurrence and also traverse terms multiple times. On the other hand explicit substitution allow us to not generate anything unnecessarily large and keep the computation at a minimum.

We consider a substitution which is lazy, meaning we use the result of a substitution, by lookup, when necessary and then proceed, but does not generate any explicit substitutions. Specifically we use Rust closures to capture Γ .

5.1.3 Normalization by Evaluation

As mentioned, we consider checking of types w.r.t. definitional equality. To do this we must have terms on normal forms, and we use the Normalization by Evaluation (NbE) for this. NbE is a process first introduced by Berger and Schwichtenberg[2] for efficient normalization of simply typed calculus, but it has since been refined for other systems in Barendregt’s lambda cube. This implementation is inspired by Christensens writeup “Checking Dependent Types with Normalization by Evaluation”[5]. The technique ties a connection between syntax and semantics.

The process of evaluating a programming language amounts to either compilation to machine code followed by execution, or by the use of an interpreter. In both cases evaluation gives meaning to said program. For instance, if the result of an interpretation is a number then the number constitutes the meaning of that program. The meaning may not be concrete but can also be functions, and since we consider typechecking that can invoke evaluation values and types live in a similar domain. Evaluation of LFSC can result in values for the built in types **type** and **kind** as well as function types such as Π types. Evaluation in general is only sensible for closed terms, but we must also consider how to handle open terms.

The process of normalization on the other hand is to transform a program into its normal form. Letting $nf(M)$ denote the normal form of term M with $\Gamma \vdash M : A$, then the following properties must hold:

- $\Gamma \vdash nf(M) : A$
- $\Gamma \vdash nf(nf(M)) = nf(M)$
- $\llbracket nf(M) \rrbracket = \llbracket M \rrbracket$

That is the normal form has the same type as the original term. The normal form is idempotent and cannot be further normalized. Finally the meaning does not change when normalizing a term. Many functions have these properties, so we further consider the normal form to be expressions which contains no redexes. A redex is a function type applied directly to an argument. Specifically the normalization is considered with respect to β -reduction. As already mentioned the process of β -reduction is slow, since it requires multiple traversals of terms. Instead we interpret the understanding of finding a normal form can as evaluation on open terms. The result of such an evaluation will not have any meaning; hence not be a value but rather a modified term with possibly unknown values. We denote these as neutral expressions. A neutral expression in general may be free variables, or application where the function point is a neutral, or, in the case of LFSC, a hole. By including neutrals, we can in fact perform evaluation on open terms. We define an evaluation reflection function $T \rightarrow \llbracket T \rrbracket$ giving meaning to terms. Then to convert the meaning back into a normal form, we define a reification function $\llbracket T \rrbracket \rightarrow T$. A normal form is then obtained by evaluation followed by reification. We describe the concrete implementation of these in 5.3.7 and 5.3.6.

5.2 Reading and formatting proofs

Because of the design decisions described above we immediately get a complication with the concrete syntax. Proofs generated by CVC5 will follow the concrete syntax, but we need a representation using De Bruijn indices. Therefore we propose an approach to loading proofs as seen in Figure 14. The proof is constructed by an external SMT solver, such as CVC5. We then have 3 programs in userspace:

1. A parser, that reads the concrete syntax.
2. A converter for translating the concrete syntax into a De Bruijn representation.
3. A transformer that will transform the De Bruijn representation into a format the kernel can read.

Then we also have another parser in kernel space, which read the appropriate format for the type checker.

For the prototyping done in this project, we do not strictly follow this pipeline, as the implementation is still user space specific, but the different components are available, except for the in-kernel parser. Ideally, data transferred from the userspace into the kernel should be in a zero-copy serializable format. From the general investigation into what crate can be compiled in the kernel, we have not found any of the major crates for zero-copy to work out of the box, such as `Cap_N_Proto` and `Rkyv`.




Figure 14: Data flow diagram of sending and processing a proof in the kernel.

5.2.1 Abstract Syntax in Rust

Despite being similar to C and C++ in syntax, Rust provides a much richer type system that allows us to create enumerations with fields, a.k.a Sum types. We might for instance define a construction for Identifiers as such:

```
pub enum Ident<Id> {
    Symbol(Id),
    DBI(u32)
}
```

An identifier can either be a Symbol if it is free or a De Bruijn index if it is bound. Terms are then defined almost identically to the abstract syntax. The major difference comes from the way we represent binders like so:

```
pub enum BinderKind {
    Pi,
    Lam,
    Let,
}
pub enum Term<Id> {
    Binder{ kind: BinderKind, var: Id,
            ty: Option<Box<Type<Id>>>,
            body: Box<Term<Id>> },
    // rest of terms
}
```

A binder is either a Π type, a λ abstraction, or a let binding. We use an option type as λ abstractions might contain an annotation but can have an anonymous type as well. Π , *let* and annotated λ all have a *Some* value for *ty*. We can reuse the same structure as terms and types are both defined by *Term*. This structure is convenient in the frontend representation of the language as this allows for simpler α -normalization. In the backend language, we split this

structure into separate constructors of the `AlphaTerm` enum. A similar structure is used for the side condition language. The `AlphaTerm` sum type looks as follows.

```
pub enum AlphaTerm<Id> {
    Number(Num),
    Hole,
    Ident(Ident<Id>),
    Pi(Box<AlphaTerm<Id>>, Box<AlphaTerm<Id>>),
    Lam(Box<AlphaTerm<Id>>),
    AnnLam(Box<AlphaTerm<Id>>, Box<AlphaTerm<Id>>),
    Asc(Box<AlphaTerm<Id>>, Box<AlphaTerm<Id>>),
    SC(AlphaTermSC<Id>, Box<AlphaTerm<Id>>),
    App(Box<AlphaTerm<Id>>, Box<AlphaTerm<Id>>),
}
```

We parameterize `AlphaTerm` by `Id` which is the data representation of symbols. In the specific implementation, we consider a `&str`, which is a reference to a fixed-sized string. We use this type over a `String` type because it is more efficient and there is no need for a term to own the string. Having terms parameterized by the `Id` type allow for easy conversion to De Bruijn levels instead of string identifiers.

5.2.2 Parsing LFSC

We use `nom` for parsing. `nom` is a parser combinator library that has evolved over the years from being mainly driven by macros, to in version 7, using composable closures. It is mainly focused on parsing bytes and hereby also `str`.

Taking the topmost function `parse_file` we structure it by composition as such:

```
pub fn parse_file(it: &str) -> IResult<&str, Vec<StrCommand>> {
    delimited(ws, many0(parse_command), eof)(it)
}
```

`delimited` takes 3 parsers and constructs a closure from it. When supplied with argument `it`, parse it with the first parser, the second, and then the third and return the result of the second parser.

We can parse term binders as such:

```
fn parse_binder(it: &str) -> IResult<&str, Term<&str>> {
    alt((
        map(
            preceded(alt((reserved("pi"), reserved("!"))),
                tuple((parse_ident, parse_term, parse_term))),
            |(var, ty, body)| binder!(pi, var : ty, body),
        ),
        ...
    ))(it)
}
```

We parse the different aspects of a binder, identifier, binding term, and the bound term, and then construct the appropriate binder.

For terms in general we must ensure the parser is robust and able to handle arbitrary nesting of parentheses. So, when parsing nonterminal terms, we first parse an open parenthesis followed

by `parse_term_`. We then try to parse a binder followed by a closed parenthesis². If we fail it parse a term or, if we can parse multiple terms, an application, as shown below.

```
pub fn parse_term(it: &str) -> IResult<&str, Term<&str>> {
    alt((
        parse_hole,
        map(parse_ident, |x| Term::Ident(Ident::Symbol(x))),
        map(parse_num, Term::Number),
        open_followed(parse_term_),
    ))(it)
}

fn parse_term_(it: &str) -> IResult<&str, Term<&str>> {
    if let res @ Ok(..) = terminated(parse_binder, closed)(it) {return res;}
    let (rest, head) = parse_term(it)?;
    let (rest, tail) = many0(parse_term)(rest)?;
    let (rest, _) = closed(rest)?;
    if tail.is_empty() {
        Ok((rest, head))
    } else {
        Ok((rest, Term::App(Box::new(head), tail)))
    }
}
```

5.2.3 Converting terms

To convert from `Term` to `AlphaTerm`, we traverse the AST and use a lookup table to update symbols appropriately. The lookup table is simply a `Vec` of names that need to be substituted. When a new binder is found we push the identifier to the end of the vector. When a symbol is met we look it up in the vector and convert it into a De Bruijn index based on its position in the vector, as seen below.

```
1 fn lookup_(vars: &[&str], var: &str) -> Option<u32> {
2     vars.iter().rev()
3         .position(|&x| x == var)
4         .map(|x| (x as u32))
5 }
```

Furthermore, we specifically map the option as follows:

```
1 pub(crate) trait Lookup<'a> {
2     fn lookup(vars: &[&'a str], var: &'a str) -> Self;
3 }
4
5 impl<'a> Lookup<'a> for StrAlphaTerm<'a> {
6     fn lookup(vars: &[&'a str], var: &'a str) -> Self {
7         lookup_(vars, var).map(|x| Ident(DBI(x)))
8             .unwrap_or(Ident(Symbol(var)))
9     }
10 }
```

One thing to note is that this approach is error-prone without careful consideration. Consider the expression: $\lambda x.((\lambda y.xy) : (\lambda z.z))$. We start by pushing x to the `vars` environment. In the body of the abstraction, we have two branches to the ascription. When transforming the term, we push y to `vars`, then we replace x with the index 1 and y with index 0. We then get to transforming the type of the ascription, and because vectors are a mutable structure when

²here binders also include `:` and `^`,

pushing `z` it will lie at `vars[2]`. After taking a branch we must thus ensure to truncate the vector to its original length. For a simple solution, we define a function `local` inspired by the effectful function `local` of the Reader monad.

```

1 fn local<'a, 'b, Input, Output>
2   (fun: impl Fn(Input, &mut Vec<&'a str>) -> Output + 'b,
3    vars: &'b mut Vec<&'a str>) -> Box<dyn FnMut(Input) -> Output + 'b> {
4   Box::new(move |term| {
5     let len = vars.len();
6     let aterm = fun(term, vars);
7     vars.truncate(len);
8     aterm
9   })
10 }
```

We create a closure that takes in a term. The closure will call `fun` with the `term` and `vars` as arguments and then it will truncate the environment to its size before `fun` was called.

We can then use the function as such:

```

1 Term::Ascription { ty, val } => {
2   let mut alpha_local = local(alpha_normalize, vars);
3   let ty = alpha_local(*ty);
4   let val = alpha_local(*val);
5   Asc(Box::new(ty), Box::new(val))
6 },
```

and hereby convert the AST to include De Bruin indices.

5.3 Typechecking LFSC

In this section, we describe the implementation of the type-checking semantics. We start by introducing the value representation obtained by inference and evaluation in Section 5.3.1. Then we describe the Signature and Contexts in Section 5.3.2 and how we handle commands in 5.3.3. We describe inference in Section 5.3.4, and type checking by definitional equality in 5.3.5. Lastly, we go over the normalization process by describing reification and evaluation in Section 5.3.6 and 5.3.7.

5.3.1 Values

Since we consider typechecking using normalization by evaluation we need a new type to describe the result of evaluation. We define them in Figure 15.

Just as `AlphaTerm`, `Value` is parameterized by an `Id` type. Now we require `Id` to implement the trait `BuiltIn`. The trait is bound by other traits and defined as such:

```
pub trait BuiltIn: PartialEq + Ord + Hash + Copy
```

The syntax means that to define `BuiltIn`, `PartialEq`, `Ord`, `Hash` and `Copy` must be implemented aswell. It must be `PartialEq` to be able to look up the `T` in the environment. It must also be `Copy`. We uses this stricter trait than `Clone`, as it allows for quick “copying” and is a satisfied criteria for both `&str` and `u32`’s that could be used for De Bruijn levels. `Hash`

```

1 pub enum Value<'term, Id: BuiltIn> {
2     Pi(bool, RT<'term, Id>, Closure<'term, Id>),
3     Lam(Closure<'term, Id>),
4     Box,
5     Star,
6     ZT,
7     Z(i32),
8     QT,
9     Q(i32, i32),
10    Neutral(RT<'term, Id>, Rc<Neutral<'term, Id>>),
11    Run(&'term AlphaTermSC<Id>, RT<'term, Id>, Rc<LocalContext<'term, Id>>),
12    Prog(Vec<RT<'term, Id>>, &'term AlphaTermSC<Id>),
13 }

```

Figure 15: Value type for LFSC³

and `Ord` is not strictly necessary but is required to use `HashMap` or `Btrees` for `Signatures`. The `BuiltIn` trait itself defines how the builtin types **type**, **int** and **rational** is defined. For `&str` this is simply a stringification of the literals, for u32 represented De bruijn levels these may be 0,1,2.

The `'term` type parameter is the lifetime of the reference to an `AlphaTerm`. We need this since many of the values have references to terms.

A value might be one of the abstractions in the term language, as these cannot be reduced further. Abstractions Π and λ contain a closure with type $RT \rightarrow \Sigma \rightarrow RT$, which when constructed closes over a local context and a term, where RT is a reference counted pointer to `Value`. The reason Σ must be passed as argument to the closure is purely a matter of the borrowing rules. If a closure is constructed with a reference to Σ then we cannot extend Σ anymore. The `Pi` value further contain its domain and a boolean value. The boolean describes the freeness of the variable in the term captured by the closure. This is extremely important for performance reasons which we describe more in detail in Section 7.1.

Values can also be one of the built in types, where

- `Box` correspond to keyword **kind**
- `Star` correspond **type**.
- and `ZT` and `QT` is **int** and **rational**.

Values can also be a value of `Z` or `Q` or a `Run`, which is simply a side condition $\{S \ M\}$.

Neutral expressions consist of an `RT` which is the type describing it, and the neutral expression it describe. The `Neutral` type can be either a neutral variable of global or local scope, a hole, or an application of a neutral term to a normal form.

```

#[derive(Debug, Clone)]
pub enum Neutral<'a, T: Copy>
{
    Var(T),
    DBI(u32),
    Hole(RefCell<Option<RT<'a, T>>>),
    App(Rc<Neutral<'a, T>>, Normal<'a, T>),
}

```

```
#[derive(Debug, Clone)]
pub struct Normal<'a, T: Copy>(pub Rc<Type<'a, T>>, pub Rc<Value<'a, T>>);
```

Lastly `Value` can be programs and run commands. Programs cannot directly be constructed by inference or evaluation, instead `Program` is used to describe the type of a side condition program, since the `Pi` constructor is insufficient. `Run` is defined for pure convenience.

5.3.2 Signature and Contexts

Signatures Σ are denoted as `GlobalContext` while Γ is used for the `LocalContext`. They have a similar interface but internally works quite differently. A global context is defined as such:

```
pub struct GlobalContext<'term, K: BuiltIn> {
    kvs: HashMap<K, TypeEntry<'term, K>>
}
```

The only field of the struct is a Hashmap with key-value pairs. We only define it as a hashmap because the implementation does not live in the kernel. For a version compatible with the kernel, we would use a `Vec` or implement a hashmap that works in the kernel. In any case, the interface is the same.

A type entry is defined as:

```
pub enum TypeEntry<'term, Key: BuiltIn>
{
    Def { ty: RT<'term, Key>, val: RT<'term, Key> },
    IsA { ty: RT<'term, Key>, },
    Val { val: RT<'term, Key> },
}
```

Notice here that this does not directly correspond to our definition in 4.2.

- The `IsA` construct corresponds directly to $x : A$, while the other two are defined purely for ease of use.
- The `Def` constructor is used for definitions stating that a constant c is a term M with type A . We mainly use this for top-level definitions as well as type inference of let bindings.
- `Val` is used for extending the environment in evaluation. This constructor can never occur in Σ but may occur in Γ . Having this constructor allow us to reuse Γ for evaluation.

The global context exposes the following functions:

```
pub fn insert(&self, key: K, ty: RT<'term, K>)
pub fn define(&self, name: K, ty: RT<'term, K>, val: RT<'term, K>)
pub fn get_value(&self, key: &K) -> ResRT<'term, K>
pub fn get_type(&self, key: &K) -> ResRT<'term, K>
```

If one tries to get the value of a `IsA` type, they get a neutral expression consisting of the stored type `ty` and a neutral symbol of the key. On the other hand if one tries to get the type of `Val` an error occurs.

The local context exposes a much similar interface but with a different underlying data structure. The local context is implemented as a linked list using reference counted pointers and much more closely represent the concatenation of Γ presented in 4.2.

```
pub enum LocalContext<'a, K: BuiltIn> {
  Nil,
  Cons(TypeEntry<'a, K>, Rlctx<'a, K>),
}
```

Most functions we use, such as `eval`, `infer` etc, needs to have access to both Σ and Γ , and thus for simplicity we define the following wrapper.

```
struct EnvWrapper<'global, 'term, T: Copy> {
  pub lctx: Rlctx<'term, T>,
  pub gctx: Rgctx<'global, 'term, T>,
  pub allow_dbi: u32,
}
```

Here `Rlctx` is a type synonym for a reference counted Γ , whereas `Rgctx` is a standard reference to Σ with lifetime `'global`.

5.3.3 Commands

Before explaining the inference algorithm we should quickly describe how commands are handled. We define a single function to handle a specific command and then apply this on an iterator of all commands presented to the proof checker. The function starts by constructing the environment wrapper, with the current Σ and an empty Γ .

For declarations we first check that the constant we want to bind $a \notin \text{dom}(\Sigma)$ and then infers the type to make sure $a : K$ or $a : A$. We then evaluate the expression and insert it, as an `ISA`, as such:

```
1 Command::Declare(id, ty) => {
2   if gctx.contains(id) {
3     return Err(TypecheckingErrors::SymbolAlreadyDefined(id))
4   }
5   env.infer_sort(ty)?;
6   gctx.insert(id, env.eval(ty)?);
7   Ok(())
8 },
```

Definitions is similarly first typechecked, but must not be of a kind level. They are then stored in the global environment as `Def` where the value is the evaluation of the term.

`Checks` is nothing more than inferring the type to check for well-typedness.

Programs are complicated for multiple reasons. As for other modifications to Σ we check that the identifier is not in $\text{dom}(\Sigma)$. We check that the return type of a program is a **type**. We then check each argument against the empty Γ and add them to another Γ' which will be used for checking the body. By this we ensure that parameters do not depend on each other. Then, before we can typecheck the body, we must first add it to Σ , since side condition programs may be recursive. Then we have to drop the `EnvWrapper` because it borrows Σ immutably and we want to borrow it mutably to insert an entry. We insert symbol `id` as a `Def` with the type as the return type of the function and a `Prog` type as value. We can then check the body to have the same type as the return type. The code very verbosely look as follows:

```

1 Command::Prog { cache: _chache, id, args, ty, body } => {
2     ... checking id is not define and return type ...
3     let mut args_ty = Vec::new();
4     let mut tmp_env = env.clone();
5     for arg in args.iter() {
6         let arg_ty = env.infer(arg)?;
7         is_type_or_datatype(arg_ty.borrow())?;
8         let ty = env.eval(arg)?;
9         tmp_env = tmp_env.update_local(ty.clone());
10        args_ty.push(ty);
11    }
12    let lctx = tmp_env.lctx.clone();
13    drop(tmp_env);
14    let typ = Rc::new(Value::Prog(args_ty.clone(), body));
15    gctx.define(id, res_ty.clone(), typ);
16
17    EnvWrapper::new(lctx, gctx).check_sc(body, res_ty)?;
18    Ok(())
19 }

```

5.3.4 Inferring Types

To infer types of LFSC we define an `infer` function for each of the constructs in the language. The functions are implemented as inherent implementations (concrete associated functions) and have the types:

```

impl<'global, 'term, T> EnvWrapper<'global, 'term, T>
where T: BuiltIn
{
    pub fn infer(&self, term: &'term AlphaTerm<T>) -> ResRT<'term, T>
    pub fn infer_sc(&self, sc: &'term AlphaTermSC<T>) -> ResRT<'term, T>
    fn infer_sideeffect(&self, sc: &'term AlphaSideEffectSC<T>)
        -> ResRT<'term, T>
    fn infer_compound(&self, sc: &'term AlphaCompoundSC<T>)
        -> ResRT<'term, T>
    fn infer_num(&self, sc: &'term AlphaNumericSC<T>) -> ResRT<'term, T>
}

```

The functions follow closely the rules in Figure 8. Taking `infer` as an example it pattern-matches on `term`, and for λ , side conditions, and holes the inference fails. And variables and symbols are simply looked up in either Σ or Γ respectively. Now for some of the more complicated rules:

inferring a `Pi` can be seen in Figure 16. We first check if the domain is a side condition. If that is the case, we infer the type of the side condition and check it to have the same type as the type of the second part of the pair. We create a Run type `val` which will be bound in the inference of the `b`. In case it is not a side condition, we simply infer the domain to have **type** and then evaluate it, to get its value. We can then update the local environment stating that `De Bruijn` index 0 in the local context `IsA { val }` type.

For application instead of interpreting multiple continuous applications as curried form, we use a flat approach in which we evaluate each argument in a loop, as presented in Figure 17. First, we infer the function point, which is saved in a mutable variable updated at each iteration of the following loop. Each argument is checked against the domain of the Π type. If the variable bound by the Π is free in the body, then we evaluate it. Otherwise, we return an arbitrary value.

```

AlphaTerm::Pi(a, b) => {
  let val =
    if let SC(t1, t2) = &**a {
      let t1_ty = self.infer_sc(t1)?;
      self.check(t2, t1_ty.clone())?;
      Rc::new(Type::Run(t1, t1_ty, self.lctx.clone()))
    } else {
      self.infer_as_type(a)?;
      self.eval(a)?
    };
  self.update_local(val).infer_sort(b)
},

```

Figure 16: Inference of Pi-term

Lastly, the body is evaluated, with either the new value or a default value added to the local environment. The result is bound to `f_ty`. This happens for each iteration and lastly, we return the bound value. In case the argument is a hole we do not check it, as it is trivially the correct type at this point. We add it to the environment and evaluate the body.

```

App(f, args) => {
  let mut f_ty = self.infer(f)?;
  for n in args {
    f_ty = if let Type::Pi(free,a,b) = f_ty.borrow() {
      if Hole == *n {
        let hole = Rc::new(Neutral::Hole(RefCell::new(None)));
        b(Rc::new(Type::Neutral(a.clone(), hole)), self.gctx)?
      } else {
        self.check(n, a.clone())?;
        let x = if *free { self.eval(n)? } else { a.clone() };
        b(x, self.gctx)?
      }
    } else {
      return Err(TypecheckingErrors::NotPi)
    }
  };
  Ok(f_ty)
}

```

Figure 17: Inference of application term

Similar inference is done for the side condition language.

5.3.5 Checking Types

We define two functions for typechecking. One takes a term `term` and a `t2` and check against it. The other takes a sideconditions as first argument, but otherwise essentially does the same. The one defined for terms can be seen in Figure 18.

We match on `term` and if it an anonymous lambda then we check *LAM*-rule of Figure 8, otherwise we infer the type of the `term` to `t1` and check `t2` and `t1` for definitional equality.

This process is two-fold and can be seen in code below. Firstly we do a value comparison between `t1` and `t2`. If not succesful we convert them to their canonical form using reification and check for equality. Generally we cannot compare values, as functions such as `Pi` has a closure inside it that cannot easily be compared. The main reason for the `ref_compare` call

```

pub fn check(&self, term: &'ctx AlphaTerm<T>, tau: RT<'ctx, T>) -> TResult<()
, T>
{
  match term {
    AlphaTerm::Lam(body) => {
      if let Value::Pi(_, a,b) = tau.borrow() {
        let val = b(mk_neutral_var_with_type(a.clone()),
          self.gctx)?;
        let env = self.update_local(a.clone());
        return env.check(body, val)
      }
      Err(TypecheckingErrors::NotPi)
    },
    _ => {
      let t = self.infer(term)?;
      self.same(t, tau)
    }
  }
}

```

Figure 18: Typechecking a term

is to fill holes. The function return a boolean of the equality between the two values, and as long as the values are one of the simple types or syntactically same neutrals it returns true. If one of the arguments is a hole it is filled with the other value. This is done using the interior mutability of Refcells. We must do it this way, as holes cannot be filled after reading back values because a hole might occur in multiple places and we must ensure that it is filled with the same value. Empirically, the `ref_compare` will return `true` most of the time. In case it returns `false`, we reify the values into their normal form and compare them.

```

pub fn convert(&self,
  t1: RT<'term, T>,
  t2: RT<'term, T>,
  tau: RT<'term, T>) -> TResult<(), T>
{
  if ref_compare(t1.clone(), t2.clone()) { return Ok(()) }
  let e1 = self.readback(tau.clone(), t1)?;
  let e2 = self.readback(tau, t2)?;
  if e1 == e2 {
    Ok(())
  } else {
    Err(TypecheckingErrors::Mismatch(e1, e2))
  }
}

```

5.3.6 Readback (Reification)

Reification or reading back semantical objects into the term language is type-directed. `read_back` takes a type and a value as arguments. First, the term to be read back is checked to be neutral, and if that is the case then we call `read_back_neutral` since `Neutral`'s `encode` their own type. Figure 19 shows the `read_back_neutral` function.

Variables can be read back directly. Holes either get read back as their inner value or as a term language hole. The application will read back the function point and the argument point, and construct an application from it. Be aware here that the argument point is a `Normal` type. It will thus call `read_back` with its type and `val`. If the value of `read_back` is not neutral we


```

fn readback_neutral(&self, neu: Rc<Neutral<'term, T>>)
    -> TResult<AlphaTerm<T>, T>
{
    match neu.borrow() {
        Neutral::DBI(i) => Ok(Ident(DBI(*i))),
        Neutral::Var(name) => Ok(Ident(Symbol(*name))),
        Neutral::Hole(hol) => {
            if let Some(ty) = &*hol.borrow() {
                self.readback_neutral(ty.clone())
            } else { Ok(Hole) }
        },
        Neutral::App(f, a) => {
            let f = self.readback_neutral(f.clone())?;
            let a = self.readback_normal(a.clone())?;
            Ok(App(Box::new(f), vec![a]))
        },
    }
}

```

Figure 19: Reification of neutral terms

pattern-match on the type. If the type is Z or Q then we can read back integers and rationals respectively. All built-in types can be read back to the built-in terms. Π types can be read back by reading back the domain, then evaluating the body and reading the body back.

5.3.7 Evaluation

We define evaluation on two levels, namely on terms and on the functional side condition language. Evaluation of the term language is straightforward. Side conditions pairs and holes cannot be evaluated. Application has a similar structure to `infer`. We consider applications in applicative order evaluation with a loop over the arguments. The function `do_app` is then called with the function and the evaluated argument:

```

pub fn do_app(&self, f: RT<'ctx, T>, arg: RT<'term, T>) -> ResRT<'term, T>
{
    match f.borrow() {
        Value::Lam(closure) => closure(arg, self.gctx),
        Value::Neutral(f, neu) => {
            if let Value::Pi(_, dom, ran) = f.borrow() {
                Ok(Rc::new(Value::Neutral(
                    Rc::new(Neutral::App(neu.clone(), Normal(dom.clone(), arg)))
                )))
            } else {
                Err(super::errors::TypecheckingErrors::NotPi)
            }
        },
        _ => Err(super::errors::TypecheckingErrors::NotPi)
    }
}

```

Functions can be either a concrete lambda abstraction in which we simply evaluate the closure, or it can be unknown. If this is the case we check that the type of the neutral value is a function type. We construct a new neutral value, where the type of the neutral value is the range of Π and the value is an application of the unknown value onto the normal expression (dom, arg) , stating that arg has type dom . For instance if we consider the application $(f \ x \ y)$, where

$f :: a \rightarrow b \rightarrow c$, then we construct:

$$Neu(b \rightarrow c, f (x : a))$$

by the first application and

$$Neu(c, (f (x : a)) (y : b))$$

after the second application. To come full circle, if we want to read back, we will get $(f x y)$ since it is already in normal form.

Evaluation of Π terms is also interesting. For standard Π constructs, we simply evaluate the domain, construct a closure around the body, check if the bound variable is free in the range and construct a $\mathbb{P}\perp$ value. If the domain of a Π is a side condition, we evaluate the side condition and the target and check for equivalence. Lastly, the result is inserted in the local context and the body is evaluated. This is what enables the execution of side conditions in the type checking.

Side conditions on the other hand may only be evaluated when part of Π type. We evaluate the side condition language by the sc function, which follows the operational semantics in Figure 12.

6 Experiments

Since the purpose of this project is to check the feasibility of an in-kernel proof-checker that can replace the eBPF verifier, we want to evaluate the implementation with proofs that resemble what would occur in a PCC context. To do so, we consider a simple verification condition generator based on the weakest precondition predicate transformers. Specifically, we consider a limited subset of eBPF consisting only of the following instructions:

$$\begin{array}{ll} \text{(Mov)} & r_d := src \\ \text{(Update)} & r_d := r_d \oplus src \\ \text{(Neg and assign)} & r_d := -src \\ & \oplus \in \{+, -, **, /, mod, xor, \&, |, \ll, \gg\} \end{array}$$

r_d denotes an arbitrary register, and src may be either, a register or a constant value of either 32 or 64 bits. Instructions can be:

- A move from src into r_d .
- A negation of a src value into r_d
- An update with one of the binary operators \oplus where $\&$ and $|$ is binary con- and disjunction, and $\ll \gg$ is logical left and right shifts.

With this, we want to do some positive testing by constructing valid programs and then check the proofs. We consider only valid programs as CVC5 does not generate proofs for satisfiable terms but rather satisfying models.

We use QuickCheck to generate arbitrary instructions with the property that register r_0 should be greater than 0 and smaller than some arbitrary value, 8192. This simulates a sequence of instructions followed by a memory access. This is a situation the eBPF verifier has been shown to be faulty at previously[15]. We ensure this property by evaluating the program by interpretation written in Haskell. If the program satisfies the property we add a prolog and epilog to the program to make a program that can be validated by the verifier. Specifically, we initialize the registers we use and include an exit command. We create the verification condition, with the postcondition that $0 \leq r_0 < n$. We convert the representation of the verification condition into SMT2 and discharge it to CVC5. If the negation of the verification condition is unsatisfiable by CVC5, then we can discharge an LFSC proof. We compare the implementation represented in this report, from now on called *lfscr* with the proof checker in C++ provided by CVC5 called *lfsc*. We use this both for finding any bugs and also to benchmark the performance of *lfscr* with a high-performance tool.

We can also check the runtime speed of loading an eBPF program into the kernel for comparison. Although this comparison is not very precise, because the loading of an eBPF program does more than just verify it, it still gives a good indication if proof checking is far more intensive than static analysis.

In creating this experiment some credit should be attributed to Ken Friis Larsen, Mads Obitsøe and Matilde Broløs. To discharge eBPF, we use Larsens <https://github.com/kfl/ebpf-tools> and to interact with eBPF we use a collection of FFI-bindings in Haskell by Obitsøe. The generation of code and evaluation is part of ongoing research by all three and me. Lastly, conversion of the verification condition to SMT2 takes heavy inspiration from Obitsøes Thesis. The VC generator is made specifically for this project.

For the benchmarking, we create programs of size 1, 10, 20, ..., 300, 400, 500, 1000. One small caveat that needs mention is that sometimes CVC5 will give a satisfiable result instead an unsatisfiable result. This is a problem somewhere in the pipeline but I have not had time to investigate this matter. For now, we settle on generating a new program of same size and try again. In the process, it showed to be unfeasible to prove the validity of programs of size 1000 in CVC5 on my i7-1165G7 CPU with 16 GB of ram, thus in the evaluation we only consider programs up to 100. This might be either the verification condition or CVC5, that is not behaving as expected, but it is out of scope for this specific project to investigate this further. Even further and more grave, the parser is not as robust as first expected, somewhere between 300-400 instructions (depending on the specific proof), the parser will stack overflow.

6.0.1 LFSC - without side conditions?

One interesting finding about the experiments is that the proofs generated from CVC5 do not include any side conditions. Therefore it may be interesting to see if a VC generator can be encoded in such a way that it produces terms with side conditions instead of purely logical connectives. If this can efficiently be done, then it may reduce both the size of the proof as well as make the type checking faster. If it is not the case, then side conditions might be removed altogether and a standard LF format would be adequate.

7 Evaluation

In this section we first represent the speed and memory usage of *lfscr* compared to *lfsc*, in Section 7.1-7.2. In this process we describe some optimizations done along the way to improve the performance. We then in Section 7.3 evaluate the completeness, correctness and discuss some general properties of the implementation.

7.1 Speed

Before we describe the results for current solution, we should briefly discuss the performance of the initial solution and some other attempts at optimizing the code, to emphasize, what small changes can do to the performance. In the first iteration of the code, the performance of *lfscr* presented in this report was atrocious. From Table 1 it should be clear that *lfscr* was extremely slow. For a single instruction, the performance is similar to *lfsc*. In this case, we essentially just check all the signatures distributed by *cvc5*[10] along with a small proof. Already for 10 instructions, my implementation was using 4 times as long to check a proof, and for a 100-instruction straight line program this difference was close to 140 times slower. We limit the result in this section to 100 instructions as anymore would have taken far too long.

instructions	1	10	100
lfsc	9.0 ± 3.3 ms	9.0 ± 3.2 ms	59.6 ± 0.7 ms
lfscr	12.7 ± 0.7 ms	36.6 ± 1.3 ms	8.3 ± 0.9 s

Table 1: First implementation lfscr Vs lfsc

Inspecting the proofs, which can be found in the repository, in the `vcgen` folder as, `benchmark_n.plf`, one can notice that there are more than 1000 local bindings for a 100-line program. This lead me to believe that using cons lists would maybe not be optimal, switching to an approach that uses `Vec` as the underlying data structure and truncating similar to the approach described in Section 5.2.3 gave a decent improvement in speed, being a couple of seconds faster than the first implementation.

instructions	1	10	100
Cons list	12.7 ± 0.7 ms	36.6 ± 1.3 ms	8.3 ± 0.9 s
Vec	18.2 ± 1.7 ms	51.8 ± 1.9 ms	6.4 ± 0.1 s

This was still early in the development and the current implementation still uses cons lists, as they provided easier implementation of the algorithm. For smaller proofs cons lists are still faster but it might be interesting to reinvestigate if using `Vec` is more efficient, now that the implementation is complete.

7.1.1 Massive speedup

Analyzing the code with `perf`, it got clear that most of the time was used in evaluating applications, namely about 60 percent of the time spent was in `eval` and `do_app`. There is nothing inherently strange about this since proofs are mainly just applications and application chains get big for larger proofs. From analyzing the *lfsc* implementation it got clear that my implementation did unnecessary computations. Considering the example from 4.6, `and_elim` is a

4 argument symbol, of which `p` is used to destruct the `holds` of the fourth argument and fill `f1`. In the example `a0 = (holds (and cvc.p (and (not cvc.p) true)))` and while the type checking that `a0 \Leftarrow holds f1` is necessary, the following call to `eval` to bind `p` in the range of the function is unnecessary since `p` does not occur free in the range. Already for this very small formula the application consists of 6 applications at the top level. This pattern appears often in LFSC proofs. Often Π types will include a parameter that does not occur free in the body, but merely exist to destruct a pattern onto an unfilled hole. So including a calculation of whether a bound variable occurs in the body and then checking the condition before evaluation can save a massive amount of computation.

This line from the application case in `infer` (along with the actual function for calculating `free`) is enough to make `lfscr` 43 times faster and relatively comparable to `lfsc`.

```
let x = if *free { self.eval(n)? } else { a.clone() };
```

Specifically we get:

instructions	1	10	50	100	200
lfsc	5.1 \pm 2.7 ms	5.7 \pm 2.7 ms	7.9 \pm 1.8 ms	59.2 \pm 2.9 ms	22.8 \pm 1.0 ms
lfscr	5.1 \pm 1.4 ms	6.9 \pm 1.0 ms	59.4 \pm 2.0 ms	193.0 \pm 4.6 ms	676.8 \pm 13 ms

Hence we now see that the `lfscr` implementation is within a 10x margin of `lfsc`. `lfsc` takes a different approach than `lfscr`. `lfsc` does everything all at once, meaning lexing/parsing and inference, and evaluation all occur in the same function in an online approach. This approach seems to reduce a lot of overhead. Especially if we consider the *perf* data, for `lfscr`, we notice that 40 percent of the time is still used `eval`. We might therefore be able reduce the execution time by embedding evaluation into the `infer` function. One thing to note is that `lfsc` also implements tail calls by `goto` statements to achieve additional performance. We may not do this in Rust, since memory management cannot easily be statically analyzed with unstructured controlflow.

From the table we can see that the result are volatile and for a more definitive answer we should consider more program points.

7.1.2 Complexity and Constants

As may be apparent from the tables introduced until now, the methodology we use may not be optimal for assessing the runtime performance, since the programs we generate vary so much in the complexity of their proofs. If we consider the graph in Figure 20, we see that the two programs follow a very similar pattern. The difference in the constant factors are still quite large and for some proofs `lfsc` is 12 times faster than `lfscr`. Figure 22 on the other hand gives a much better view into the volatility of `lfscr` in which the running time of `lfscr` is only twice that of `lfsc`. This suggests that there might be other optimization points, similar to the one described in the previous section, to eliminate unnecessary computations. Figure 21 suggests that we have a complexity issue in `lfscr` but because of the volatility it is difficult to draw a conclusion from these data points. Here it is very unfortunate that the parser cannot handle more large enough programs. In any case, the `lfsc` implementation seems to be linear, since it has a ms/number of instructions ratio that is consistently < 1 , as seen in Figure 22.




Figure 20: Runtime of *lfsc* Vs *lfscr*, logarithmic scale




Figure 21: Runtime of *lfsc* Vs *lfscr*, linear scale




Figure 22: Runtime of *lfsc* Vs *lfscr*, ms/number of instructions

ADDENDUM These benchmarks were done before I realized that *lfscc* can be built in both a debug and release version. In the release version, it is consistently 2-3 times faster than the results presented here. This suggests that a proof checker can indeed be efficiently implemented, but the approach used in this project is not ideal. It might be possible to reduce the overhead by quite a bit, but it is unlikely that we can reach exactly the same level of performance of *lfscc* with a staged process like done in *lfsr*.

7.1.3 Formal checking vs static analysis

We should not only consider the execution time of *lfsr* in terms of other implementations. We should also compare the runtime with how long the verifier runs. It is not immediately as easy to benchmark the performance of kernel functions, although we could potentially have used eBPF to benchmark the verifier. Instead, we settle for a simpler but more inaccurate solution, where we benchmark the entire loading call. With this, we get the following running times:

Program size	1	10	100	1000
Loading time of eBPF	$57.3 \pm 8.3 \mu s$	$58.1 \pm 3.2 \mu s$	$134.3 \pm 2.9 \mu s$	$1.6 \text{ ms} \pm 144.9 \mu s$

It should here be clear that the verifier is a lot faster. Even a 100-line program only takes 134 nano-seconds, which percentage-wise is significantly faster than checking a proof. Instead of directly comparing the running times of formal checking vs static analysis, we should instead consider them from a pragmatic perspective. The question then becomes, is it worth spending a second or two, to load a program that is guaranteed to not be malicious, or is it more worth to be able to load programs extremely fast?

7.2 Memory

We should consider the memory usage of the implementation in two manners.

First, the size of proofs plays a key role in the feasibility of using proof-carrying code. A proof for a single instruction program (actually 4 with pre-initialization and the epilog), is 2.7KB in size, while 10 instructions are 8.6KB and 100 instructions are 109KB. For larger programs such as 400 and 500 lines, the size is 668KB and 706KB. So the proofs, at least for straight-line programs, scale linearly (or close) with roughly 1-2 KB per instruction. Encoding the proofs in a more compact binary format could make these sizes even smaller. The sizes in themselves are not alarming and could still see use in devices with limited memory. For embedded devices the filesizes may be too big, but these devices will probably not run Linux.

Secondly, we should also look at how much memory the type checker uses. Running both *lfsr* with the 1,10 and 100 line proofs, we get the following memory usage:

Program size	1	10	100	200	300
peak memory	1.3MB	1.8MB	5.7MB	21.1MB	23.6MB
peak RSS	9MB	15.7MB	25.3MB	47.7MB	51.6MB
temporary allocations:	50.13 %	46 %	40 %	40.7 %	40.2 %

From these results, we see that *lfscr* does not use a massive amount of memory, for smaller programs. For a 100-line eBPF program we use at most 5.7 MB at a time. For larger programs, we allocate 23.6 MB for a 300-line program, and for the entire lifetime of the program use 51.6MB.⁴ What is most interesting is that 40 % of allocations are temporary and for smaller programs even higher. This suggests that we do some unnecessary computations and that we maybe should use another approach than reference counted pointers. This especially becomes noticeable, when similar diagnostics are done for *lfscc*. For the 300-line program only 6MB of memory is used at its peak, while it uses 14.2MB overall and only 6% of allocations are temporary. One thing to keep in mind is that about 1/3 of allocations are leaked. This is not ideal, but for very shortlived programs such as *lfscc* it is not a big deal. On the other hand for a program that runs in the kernel memory leaks are problematic.

In any case, we can again see that we can check large proofs without many resources needed. But that an “all in one” solution presented by *lfscc* could be worth prototyping in either pure C or in Rust.

7.3 LFSCR - strong suits and weaknesses

As the previous section described the performance of *lfscc* suggests that a more efficient approach exists, than we have at the moment. This implementation does have a couple of features that are worth taking into consideration as well. It is implemented completely in safe Rust, meaning we cannot have any illegal memory access. We have further ensured the implementation to be panic-free. There is however a problem with the parser at the moment, where it stack overflows for very large nested applications. This is unacceptable in a kernel context and should be fixed. This might be the most desirable property for a program that is designed to run inside the kernel, as “proofs” could exploit such a vulnerability.

Equally an implementation should be robust in the amount of time it takes to check the proof. We showed before the performance difference in checking if the occurrence of a variable was free could improve the performance by 43 times. This immediately shows that we should also consider some sort of time limit for how long a proof must be, since a malicious “proof” could slow down a system massively.

lfscr has an additional advantage over *lfscc* when considering the position in a PCC architecture. Checking the proof has not been tampered with is straightforward and already implemented unintentionally. In its current state, the LFSC proofs discharged from CVC5 always contain the following pattern:

```
... POTENTIAL BINDINGS ...
(# a0 (holds x)
(: (holds false)
... ACTUAL PROOF...
```

here *x* is the formula unsatisfied by CVC5. Given that an in-kernel VC generator constructs its verification condition as a *AlphaTerm*, then the check is nothing more than normalizing the verification condition and the *a0* of the proof and check for equality.

⁴Note that this memory also includes some heaptrack overhead.

The experiment has not only provided useful insight into the performance of the implementation; it also establishes confidence that the proof checker works as expected and follows the semantics from 4.3. Checking the signatures along with the generated proofs suggests that mostly all parts of the type checker are correct. All matters of the term language are covered, and most of the side condition language is also checked. At the moment `markvar` and `if_marked` are left incomplete. The main reason for this is that there are currently no signatures distributed by CVC5 that include them. The side condition language could be tested more thoroughly as only a single larger test has been conducted by the $P \wedge \neg P$ unsatisfiability proof from 4.6. Despite the example being rather small, it tests a large part of the side condition language, both constant and program application, match constructs, branching, and numerical functions. One point where the implementation is inherently wrong is the usage of `i32s` for the representation of integers and rationals these should be unbounded integers. This is not a problem for bit-vector proofs, but only for arithmetic logic. I have however left the representation as is for now, as I have not been able to find a library that efficiently implements unbounded integers and rationals and are compatible with the kernel requirements.

Albeit the implementation does not run in the kernel, the implementation only uses the `core` and `alloc crate` along with `nom`, which I have been successful in compiling and simple examples of in a kernel module. Hence there is nothing theoretical stopping us from compiling *lfscr* into the kernel. The major work that should be done here is to make every allocation fallible by using the `try_new` counterparts to `new` allocations and implementing a simple `From trait` to easily convert allocation errors into type-checking errors.

8 Is PCC a good idea?

Even with a Rust implementation that promises memory safety and has no unexpected errors that can crash the program, the answer is not definite at this time. It might still not be feasible to use LFSC for an in-kernel proof checker as part of a larger proof-carrying code architecture, since a lot of questions are still unanswered. The eBPF verifier does a lot more than just validate instructions of a bytecode format. It checks validity in memory alignment, user-rights, does program rewrites, and much more. Some of these can be encoded into a proof, but others may be harder to realize. Especially user-rights can prove as a challenge, since it either requires the code producer to make the proof themselves, meaning eBPF programs are not that easily distributed over machines, or they have to be patched in some way. Another possibility is only checking capabilities as a separate stage before checking the proof, but this may reduce some functionality of some “features” of eBPF in its current form for some users. Thus there is still a lot more work to be done in the architectural construction of a PCC system.

Another pressing matter is the execution time. We have seen that proof-checking of validity in eBPF programs (at least straight-line programs) can be efficiently done, but with this implementation, we are not quite there yet. The benchmarking showed promises in a few different places. Using `Vec` instead of `cons` lists may be useful for larger proofs, and it would further be interesting to investigate if a modified version of the data layout would prove useful. For instance, we might be able to use tagged pointers or at least a more compact data format for `Value`’s and `Neutral`’s to make the program both more memory and runtime efficient. Furthermore, the benchmarking we have done may not be entirely appropriate for determining the feasibility as we have only included straight-line programs and no control flow constructs. In the end, this can make proofs more complicated.

Despite all this, the implementation we present is rather small and consists of only 2400 lines of code compared to 19000 in the verifier. Bugs are hence less likely to appear. In any case, we do not completely discard the idea of PCC in the kernel as it does show promises and with time could be a generally decent replacement for the eBPF verifier.

9 Future work

We have presented only a small part of a proof-carrying code architecture, so naturally for a final conclusion on the matter, discussed in this work, an in-kernel verification condition generator should be defined. We suggest such a proof checker represents the logic in the form of LFSC.

The implementation discussed should be modified to run inside the kernel. Because of the preliminary work, this should be a minor task.

More work should be done in making the implementation fast.

Lastly, a more comprehensive analysis and design between PCC and the eBPF verifier should be conducted. Specifically, which parts of the verifier should co-exist with PCC. One such example is checking of user-rights.

10 Conclusion

We have in this report given a brief overview of how the eBPF verifier works, and we have discussed some concerns with the current eBPF verifier. This motivates the intent of making a proof carrying code architecture as part of the Linux kernel. In this project we have

- Described how we can use the reasonably new feature of Rust inside the kernel, to make a PCC architecture.
- Described the dependently typed language LFSC, which uses the Curry-Howard isomorphism to construct logics.
- Described and implemented, although not completely, a typechecker for the LFSC language. The typechecker in its current form is not compileable inside of the kernel but uses only features of the Rust language, with some slight modifications, such as enabling the reference-counted smart-pointer. The implementation of the typechecker is reasonably simple in structure, with only 2400 lines of code and uses purely safe features in Rust and thus provide the necessary safety to run inside the kernel.
- Done experiments that generates arbitrary straight line eBPF programs and performs a naive verification condition on them to generate LFSC proofs with CVC5. We used these experiments to ensure the correctness of the implementation as well as providing a framework for benchmarking the performance.
- Found that for verification conditions for straight-line programs consists of no side conditions, which suggest they might be unnecessary.

- By accident also facilitated the foundation for ensuring proofs correspond to the program it is supplied with in a PCC setting. The argumentation here is that the proof encodes the verification condition as an LFSC term, and if the in-kernel proof checker generates verification conditions in LFSC, then checking the requirement is merely normalization followed by an equality check.

In this process we have found the implementation described in this report to be inferior in performance to the *lfsc* proof checker from CVC5, by up to 12 times slower execution time, which suggest a modular approach using normalization by evaluation is not the ideal approach. Although the execution time is higher for this implementation it is still within reason and may, with some optimizations, be decent enough for an in kernel verifier.

Thus, while we cannot say anything definite at this point, we have provided a basis for future work to make a PCC architecture in the Linux kernel.

References

- [1] *Adding syscalls*. URL: <https://www.kernel.org/doc/html/v4.10/process/adding-syscalls.html> (visited on 05/29/2023).
- [2] U. Berger and H. Schwichtenberg. “An inverse of the evaluation functional for typed lambda-calculus”. In: *[1991] Proceedings Sixth Annual IEEE Symposium on Logic in Computer Science*. 1991, pp. 203–211. DOI: 10.1109/LICS.1991.151645.
- [3] *BPF Compiler Collection (BCC)*. URL: <https://github.com/iovisor/bcc> (visited on 05/26/2023).
- [4] *Carcara*. URL: <https://github.com/ufmg-smite/carcara> (visited on 05/29/2023).
- [5] David Thrane Christiansen. *Checking Dependent Types with Normalization by Evaluation: A Tutorial (Haskell Version)*. 2019.
- [6] Edsger W. Dijkstra. “Guarded Commands, Nondeterminacy and Formal Derivation of Programs”. In: *Commun. ACM* 18.8 (August 1975), pp. 453–457. ISSN: 0001-0782. DOI: 10.1145/360933.360975. URL: <https://doi.org/10.1145/360933.360975>.
- [7] *eBPF Helper calls*. URL: <https://ebpf.io/what-is-ebpf/#helper-calls> (visited on 05/26/2023).
- [8] *eBPF verifier*. URL: <https://docs.kernel.org/bpf/verifier.html> (visited on 05/29/2023).
- [9] Robert Harper, Furio Honsell, and Gordon Plotkin. “A Framework for Defining Logics”. In: *J. ACM* 40.1 (January 1993), pp. 143–184. ISSN: 0004-5411. DOI: 10.1145/138027.138060. URL: <https://doi.org/10.1145/138027.138060>.
- [10] *LFSC signatures*. URL: <https://github.com/cvc5/cvc5/tree/main/proofs/lfsc/signatures> (visited on 05/28/2023).
- [11] *LFSC: A High-Performance LFSC Proof Checker*. URL: <https://github.com/cvc5/LFSC> (visited on 05/29/2023).
- [12] *libbpf*. URL: <https://github.com/libbpf/libbpf> (visited on 05/26/2023).

- [13] *Loading process of eBPF programs*. URL: <https://ebpf.io/static/7eec5ccd8f6fbaf055256da491b5f15/loader.png> (visited on 05/26/2023).
- [14] George C. Necula. “Proof-Carrying Code”. In: *Proceedings of the 24th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. POPL ’97. Paris, France: Association for Computing Machinery, 1997, pp. 106–119. ISBN: 0897918533. DOI: 10.1145/263699.263712. URL: <https://doi.org/10.1145/263699.263712>.
- [15] Manfred Paul. *CVE-2020-8835: Linux Kernel Privilege Escalation via Improper eBPF Program Verification*. 2020. URL: <https://www.zerodayinitiative.com/blog/2020/4/8/cve-2020-8835-linux-kernel-privilege-escalation-via-improper-ebpf-program-verification> (visited on 05/25/2023).
- [16] Simon Scannell. *Fuzzing for eBPF JIT bugs in the Linux kernel*. 2021. URL: <https://scannell.io/posts/ebpf-fuzzing/> (visited on 05/23/2023).
- [17] Aaron Stump et al. “SMT Proof Checking Using a Logical Framework”. In: *Formal Methods in System Design* 42 (February 2013). DOI: 10.1007/s10703-012-0163-3.